

## El caso para Base de Datos NO relacionales

- Grandes volúmenes de datos y sistemas basados en la Web
- Escalabilidad, elasticidad, eficiencia, tolerancia a fallas
- Replicación y particionado de datos.
- No existen Joins. El costo de joins en bases de datos relacionales llega a ser demasiado elevado cuando se trata de tablas con grandes cantidades de tuplas.
- Se aplica DES-normalización
- Clusters de procesadores y centros de datos
- Conceptos de paralelismo y sistemas distribuidos embebidos en el modelo de datos y el despliegue físico del modelo de datos.
- Diseño del modelo de datos basado en las CONSULTAS que debe soportar la aplicación, las tablas se diseñan para procesarlas de manera eficiente.

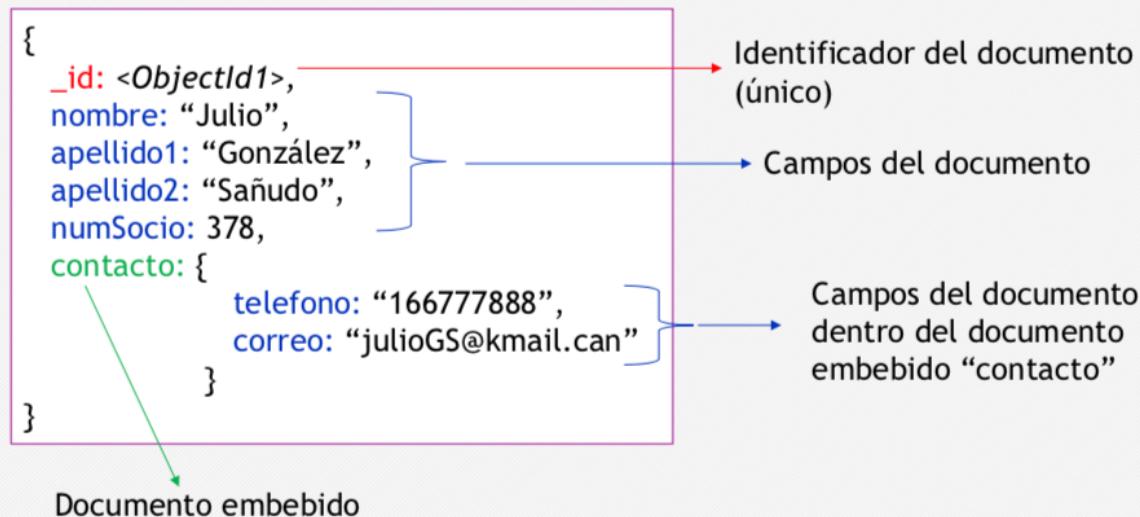
NoSQL = Not Only SQL databases

Existen al menos tres tipos de bases de datos NoSQL.

- Pares **llave-valor**, donde “valor” puede ser cualquier objeto, binario, JSON, se utilizan en situaciones donde es necesario almacenar y recuperar de manera eficiente objetos individuales tales como imágenes.
- **Documentales**, los objetos tienen mayor estructura, XML, JSON, texto, son documentos que pueden tener una estructura de árbol y/o pueden referenciar a otros documentos. Por ejemplo un documento que describe una orden de compra que contiene dirección de entrega del pedido, dirección de envío de la factura, y la lista de productos con la cantidad requerida de cada producto (en bases de datos relacionales se requieren varias tablas).
- **Columnares**, tienen la forma de tablas donde se especifican familias de columnas, muy flexibles, y se proporcionan operaciones que permitan ejecutar funciones sobre las familias de columnas. No tienen las restricciones de tablas con campos fijos del modelo relacional. Las filas pueden contener diferentes subconjuntos de columnas, y dentro de cada familia de columnas pueden haber varias filas. Explotan localidad de accesos en grandes volúmenes de datos.
- **Grafos**, permiten almacenar y recuperar eficientemente relaciones entre pares de datos “Objeto → Relación → Objeto”.

- MongoDB es una Base de Datos NoSQL orientada a documentos en formato BSON.
  - Similar a JSON, con algunas particularidades:
    - Tipos de datos adicionales
    - Diseñado para ser más eficiente en espacio.
      - Aunque en ocasiones un JSON puede ocupar menos que un BSON.
    - Diseñado para que las búsquedas sean más eficientes.

- Ejemplo de documento:



- Los datos son accedidos mediante operaciones *CRUD* (*Create, Read, Update, Delete*):

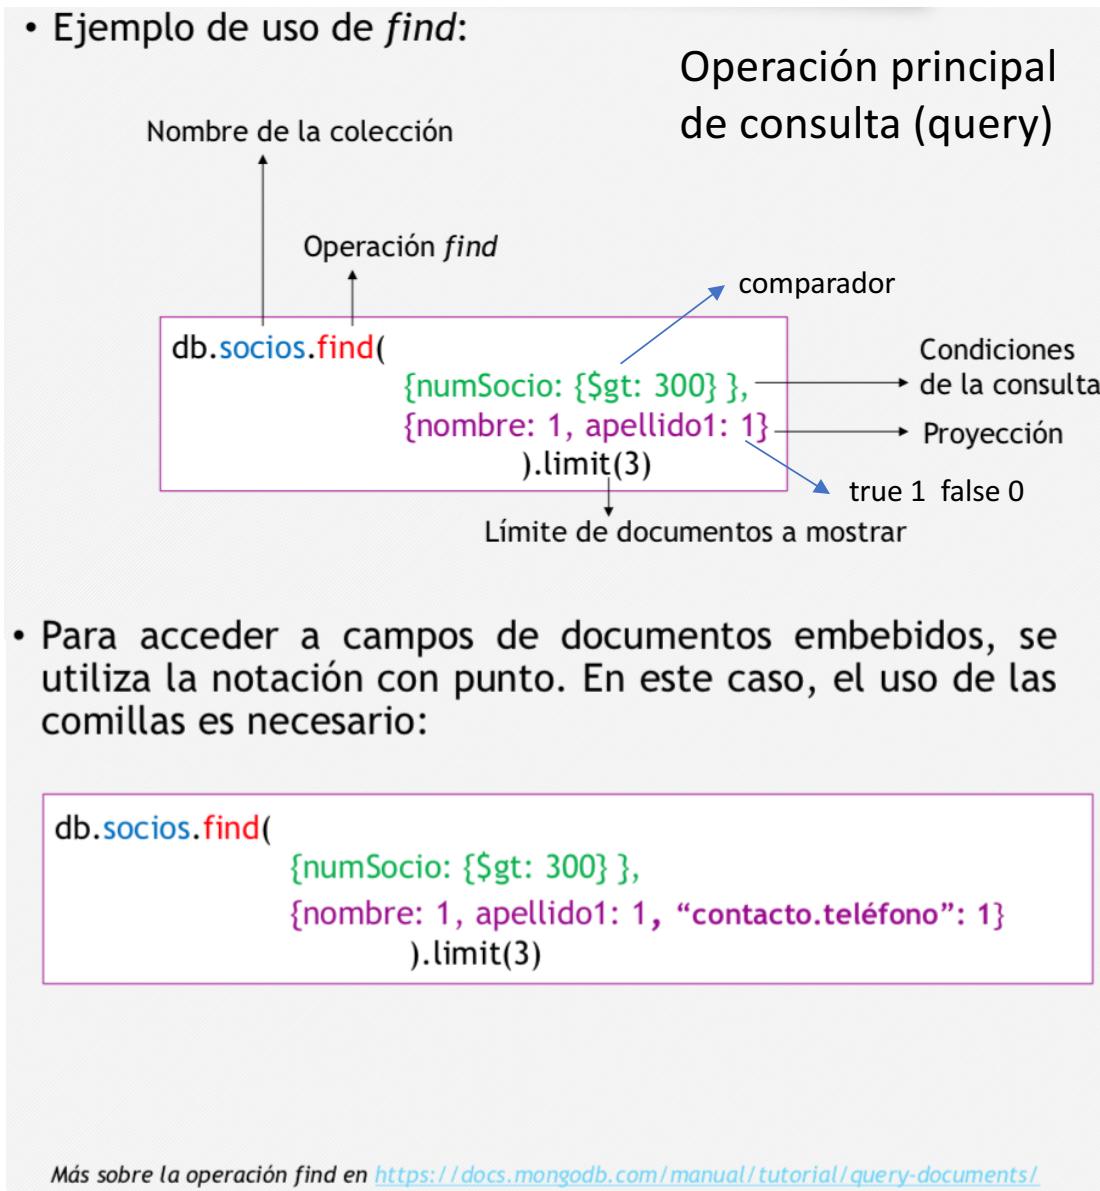
- Find: lectura (Read) de datos. Símil con el SELECT en SQL.
  - Insert: escritura (Create) de datos.
  - Update: actualización de valores en los datos.
  - Remove: borrado (Delete) de datos.
- La ejecución de cualquiera de estas instrucciones sigue el siguiente esquema:

```
db.<nombre_colección>.<nombre_operación>(<condiciones de la operación>)
```

- Ejemplo de colección:

```
{  
  _id: <ObjectId1>,  
  nombre: {  
    apellido1: <ObjectId2>  
    nombre: {  
      _id: <ObjectId3>,  
      nombre: "Julio",  
      apellido1: "González",  
      apellido2: "Sañudo",  
      numSocio: 378,  
      contacto: {  
        telefono: "166777888",  
        correo: "julioGS@kmail.can"  
      }  
    }  
  }  
}
```

- Ejemplo de uso de *find*:



SELECT item FROM inventory WHERE status="A"	db.inventory.find( {status:"A"}, {item:1,_id:0})
SELECT item FROM inventory WHERE status <>"A"	db.inventory.find( {status:{ \$ne:"A"}}, {item:1,_id:0})
SELECT * FROM inventory WHERE status="A" or item = "paper"	db.inventory.find( {\$or:[{status:"A"}, {item:"paper"}]})
SELECT * FROM inventory WHERE status="A" ORDER BY item desc	db.inventory.find( {status:"A"}).sort( {item:-1})
SELECT count(*) FROM inventory WHERE status="A"	db.inventory.count( {status:"A"})  db.inventory.find( {status:"A"}).count()

- Existen diferentes tipos de operaciones:

De igualdad (\$eq):

```
db.socios.find({numSocio: {$eq: 300}})
```

*Forma implícita de la operación \$eq*

```
db.socios.find({numSocio: 300})
```

Mayor que (\$gt):

```
db.socios.find({numSocio: {$gt: 300}})
```

Mayor o igual que (\$gte):

```
db.socios.find({numSocio: {$gte: 300}})
```

Menor que (\$lt):

```
db.socios.find({numSocio: {$lt: 300}})
```

Menor o igual que (\$lte):

```
db.socios.find({numSocio: {$lte: 300}})
```

Diferente a (\$ne):

```
db.socios.find({numSocio: {$ne: 300}})
```

Pertenece (\$in) o no (\$nin) a un conjunto

```
db.socios.find({numSocio: {$in: [300, 301]}})
```

```
db.socios.find({numSocio: {$nin: [300, 301]}})
```

- Operadores lógicos:

Ambas condiciones han de cumplirse(\$and):

```
db.socios.find({$and: [{nombre: "Juan"}, {apellido1: "Galindo"}]})
```

Alguna de las condiciones ha de cumplirse (\$or):

```
db.socios.find({$or: [{nombre: "Juan"}, {apellido1: "Galindo"}]})
```

No debe de cumplir la condición (\$not):

```
db.socios.find({numSocio: { $not: {$eq: 300} }})
```

Que no se cumpla ninguna condición (\$nor):

```
db.socios.find({$nor: [{nombre: "Juan"}, {apellido1: "Galindo"}]})
```

- *Sort*: ordena los resultados según los campos dados. Un valor de 1 indica ordenamiento ascendente, y un valor de -1 ordenamiento descendente:

```
db.socios.find({numSocio: {$eq: 300}}, { nombre: 1, apellido1: 1 }).sort({nombre: 1})
```

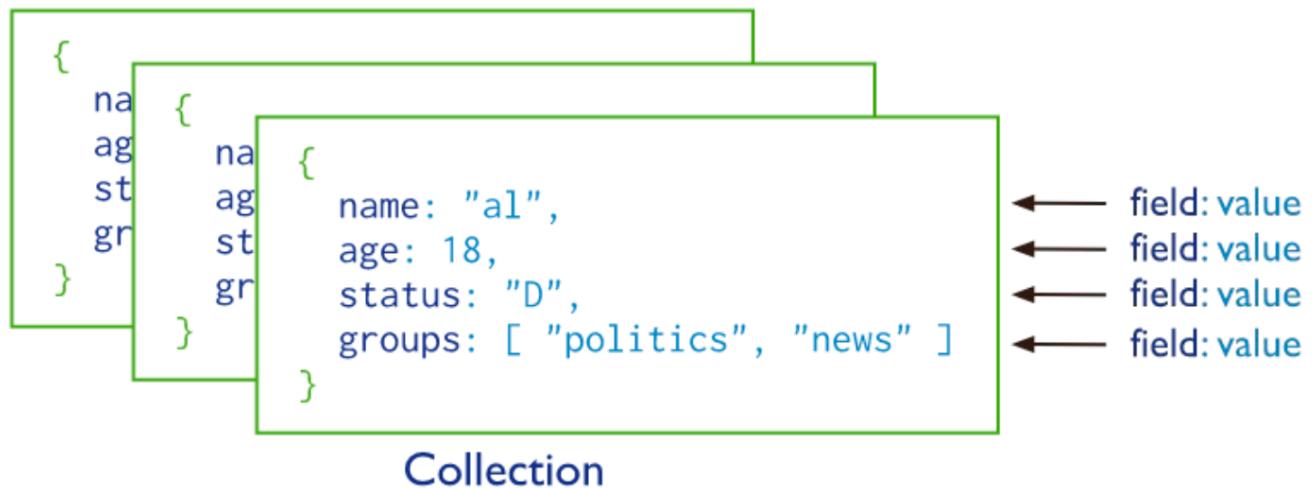
- *Count*: cuenta el número de veces que aparece en los documentos de una colección un valor en un campo.
- Ejemplo: suponiendo que en una colección llamada “series” tenemos un campo puntuación, y los siguientes documentos:

```
{puntuacion: 5}, {puntuacion: 3}, {puntuacion: 5}, {puntuacion: 5}, {puntuacion: 2},
```

La siguiente instrucción retornaría 3 (3 veces que se repite el valor 5 en el campo puntuación):

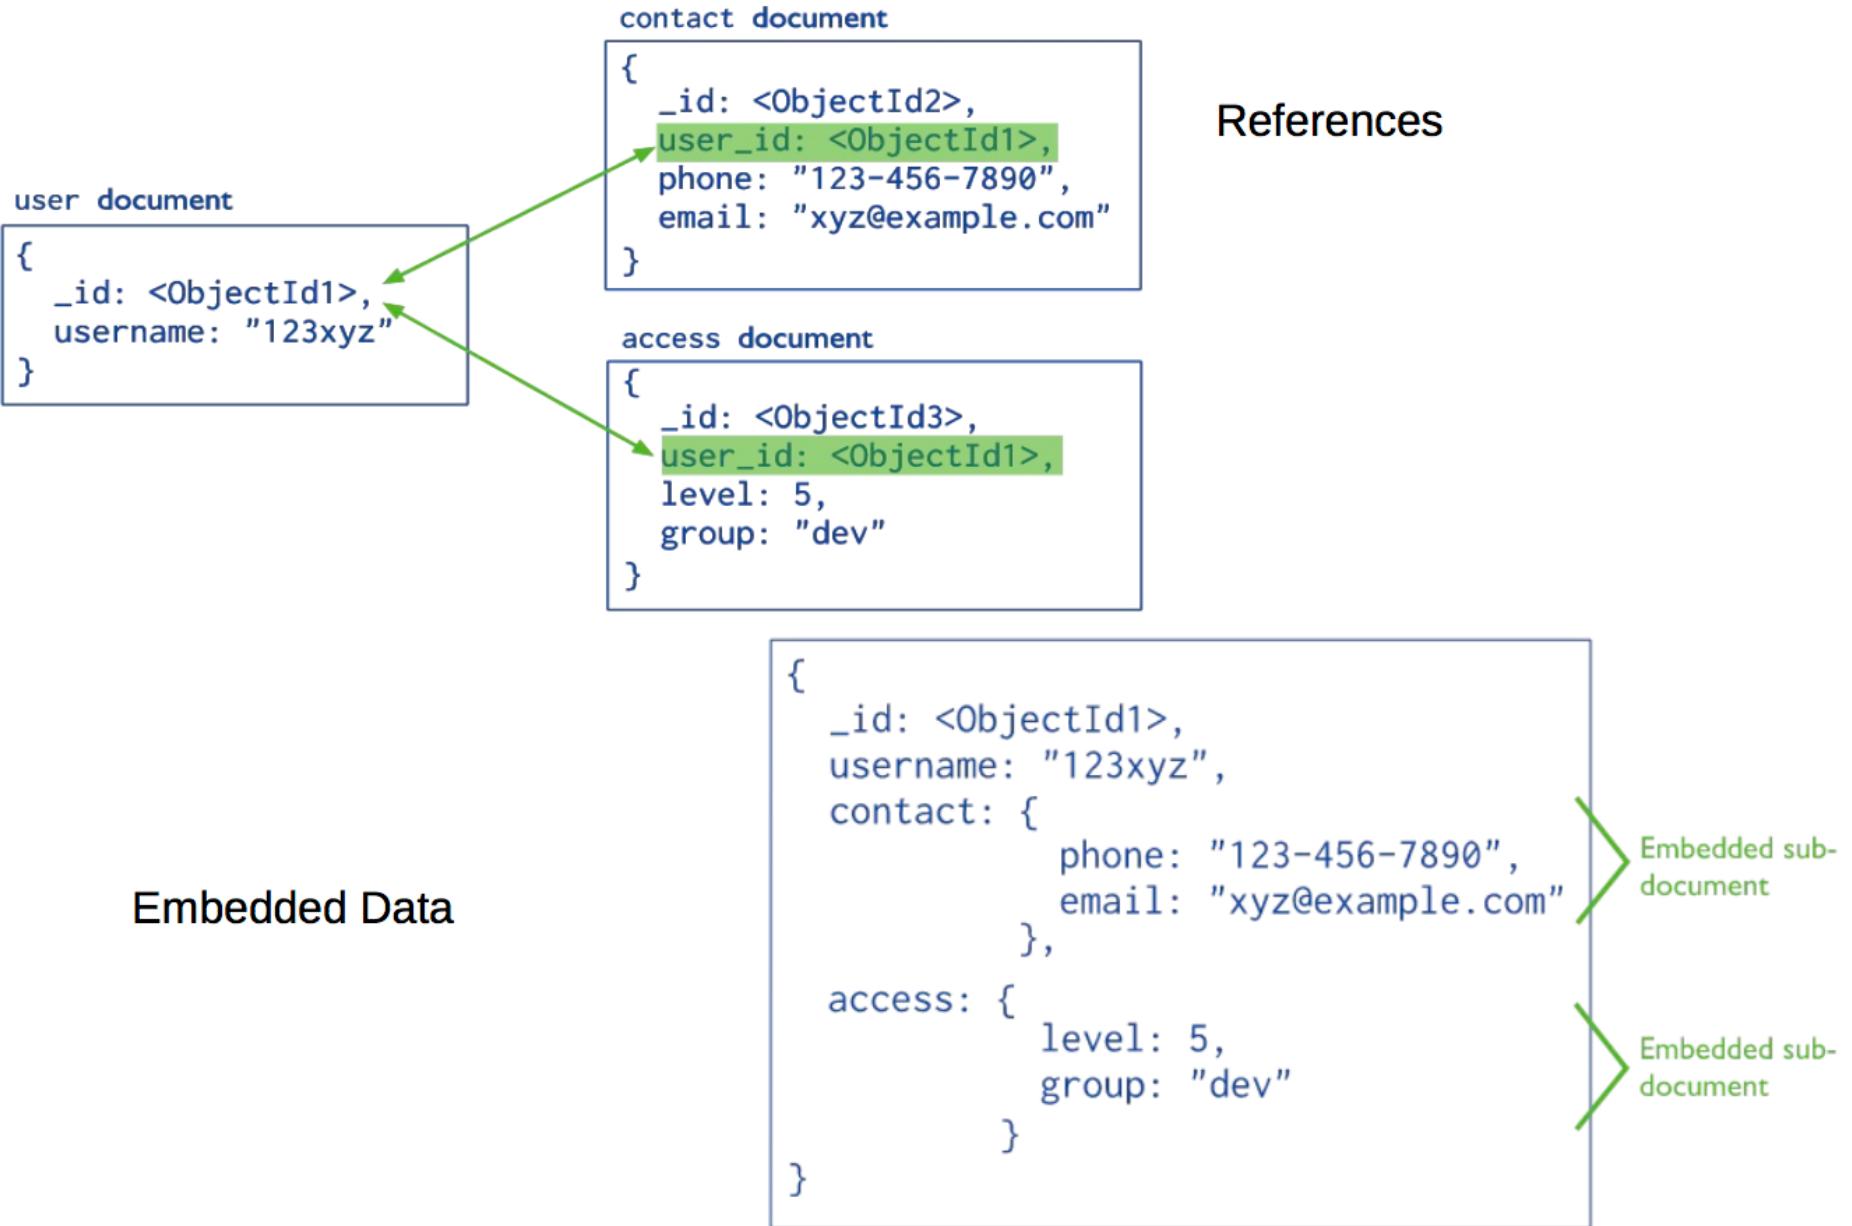
```
db.series.count({puntuacion: 5})
```

## DOCUMENT AND COLLECTION

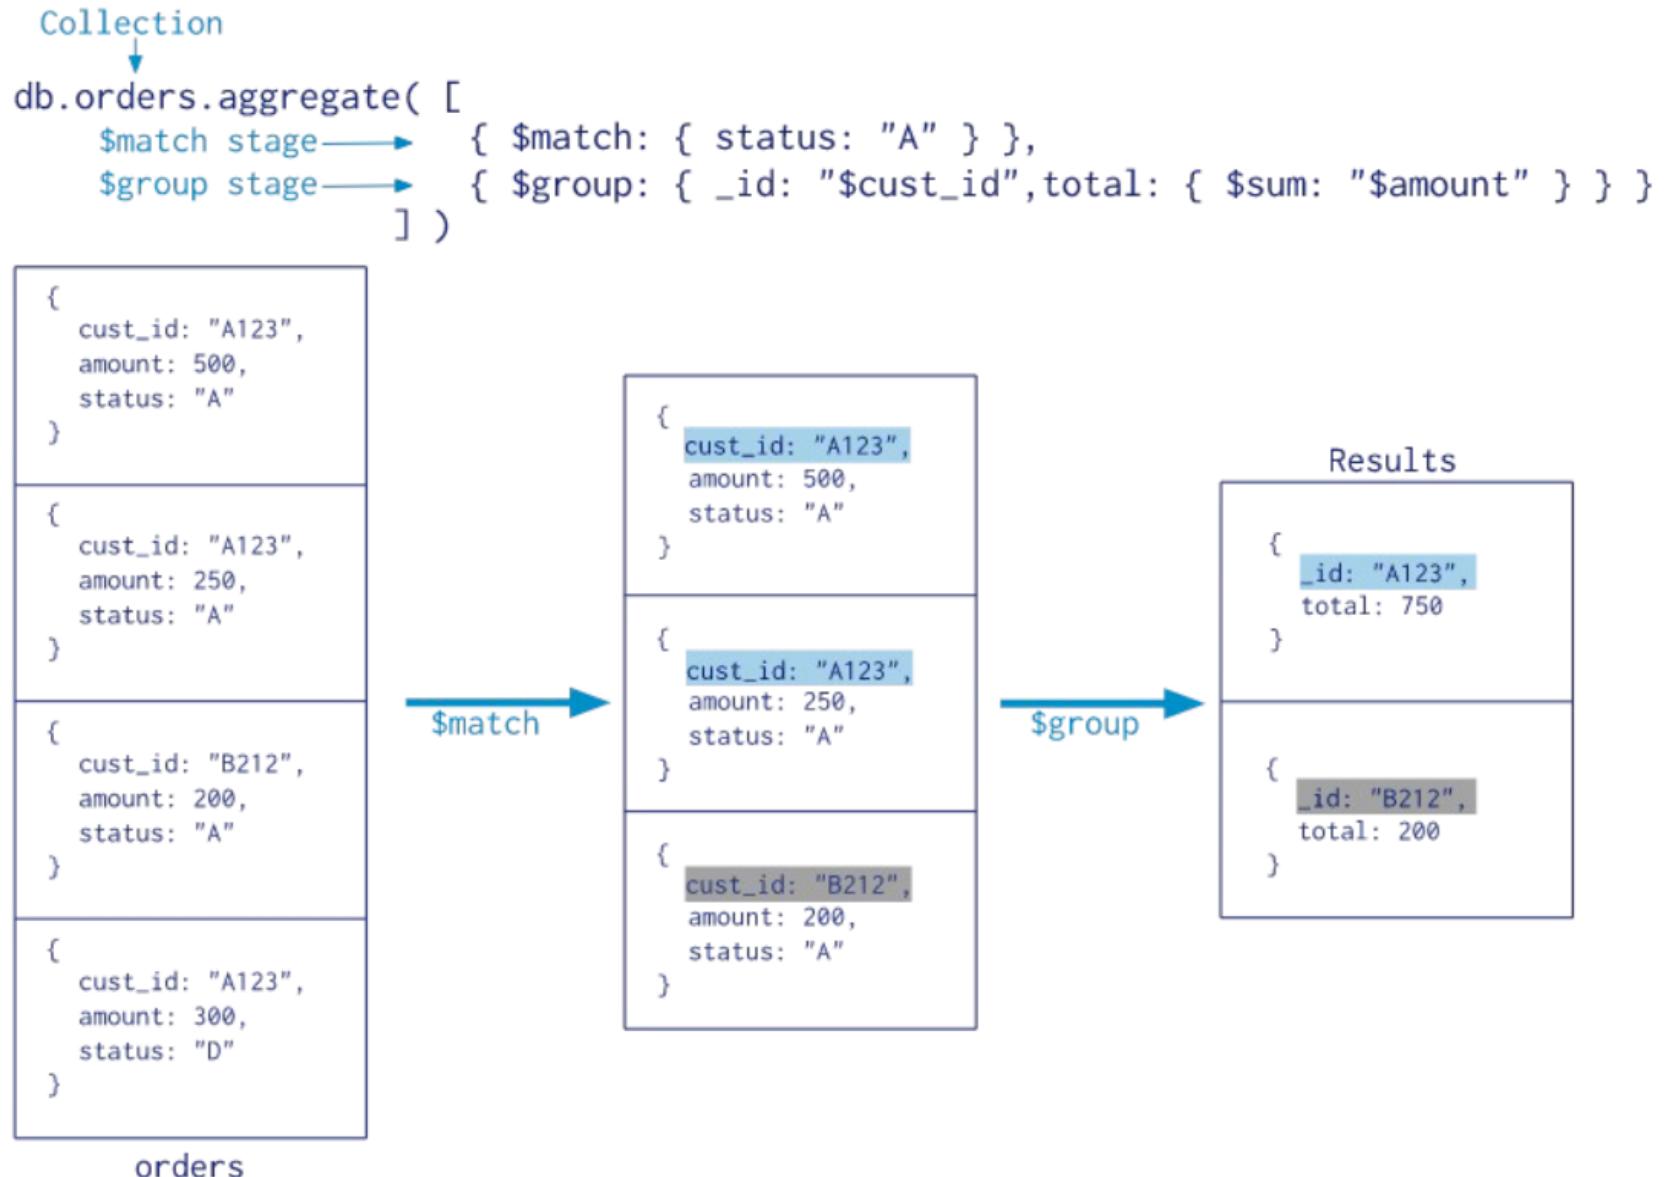


```
var mydoc = {  
  _id: ObjectId("5099803df3f4948bd2f98391"),  
  name: { first: "Alan", last: "Turing" },  
  birth: new Date('Jun 23, 1912'),  
  death: new Date('Jun 07, 1954'),  
  contribs: [ "Turing machine", "Turing test", "Turingery" ],  
  views : NumberLong(1250000)  
}
```

## DOCUMENT

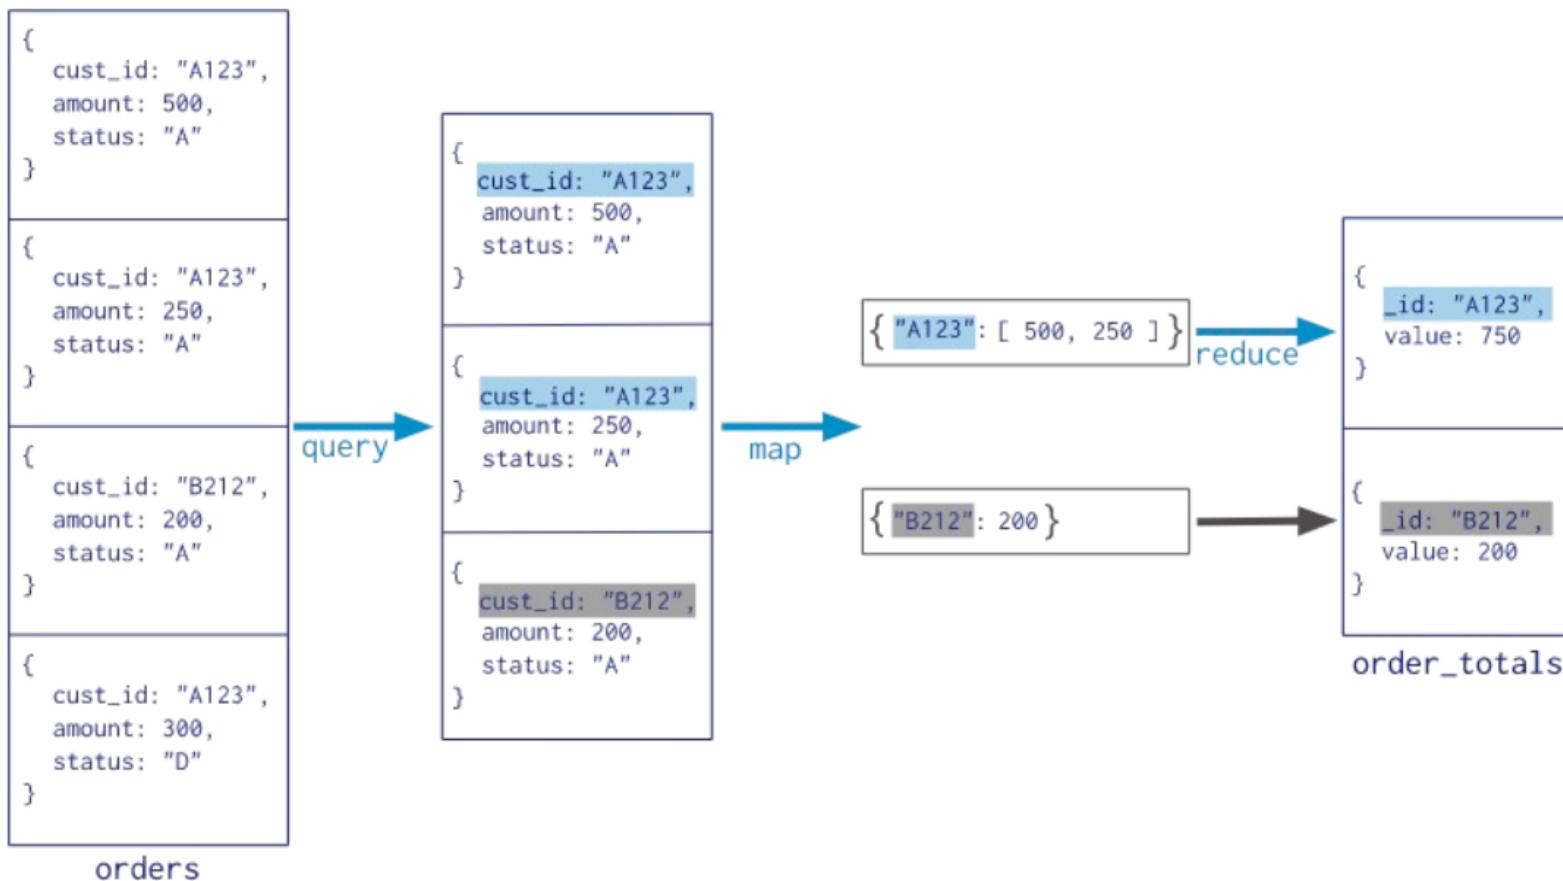


# AGGREGATION PIPELINE

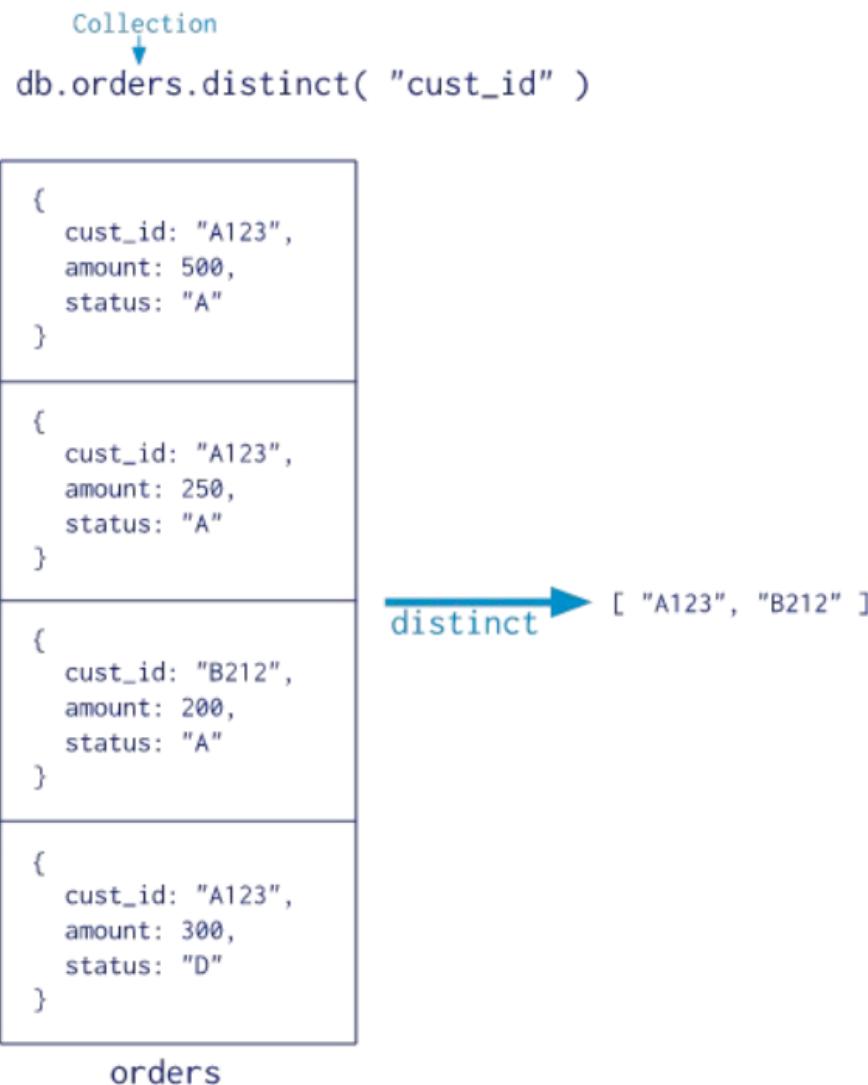


## AGGREGATION MAP-REDUCE

```
Collection  
↓  
db.orders.mapReduce(  
    map → function() { emit( this.cust_id, this.amount ); },  
    reduce → function(key, values) { return Array.sum( values ) },  
    {  
        query → { status: "A" },  
        output → "order_totals"  
    }  
)
```



## AGGREGATION SINGLE PURPOSE



## SEARCH TEXT

```
db.stores.insert(  
[  
    { _id: 1, name: "Java Hut", description: "Coffee and cakes" },  
    { _id: 2, name: "Burger Buns", description: "Gourmet hamburgers" },  
    { _id: 3, name: "Coffee Shop", description: "Just coffee" },  
    { _id: 4, name: "Clothes Clothes Clothes", description: "Discount clothing" },  
    { _id: 5, name: "Java Shopping", description: "Indonesian goods" }  
]  
)
```

```
db.stores.createIndex( { name: "text", description: "text" } )
```

```
db.stores.find( { $text: { $search: "java coffee shop" } } )
```

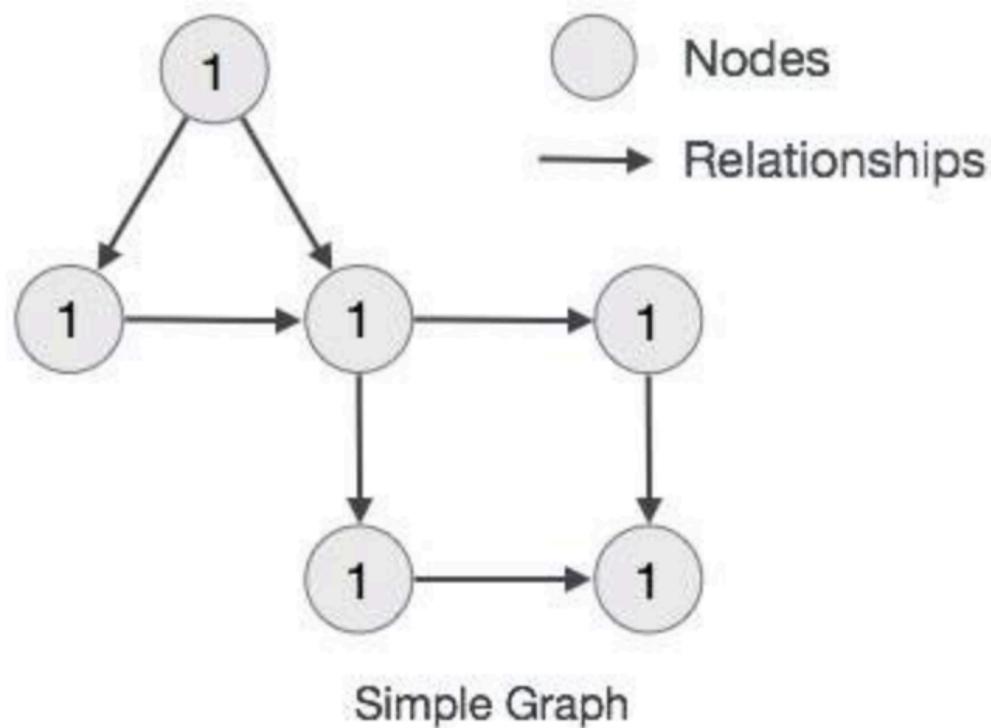
```
db.stores.find( { $text: { $search: "java \"coffee shop\"" } } )
```

```
db.stores.find( { $text: { $search: "java shop -coffee" } } )
```

# Gestor de Bases de Datos en Grafos Neo4J

<http://www.tutorialspoint.com/neo4j>

Simple Property Graph example

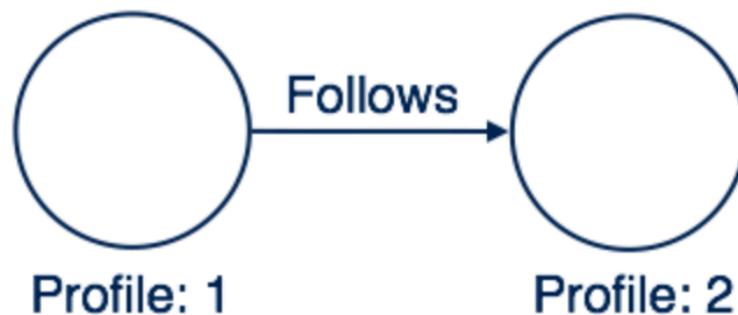




**Profile Node**

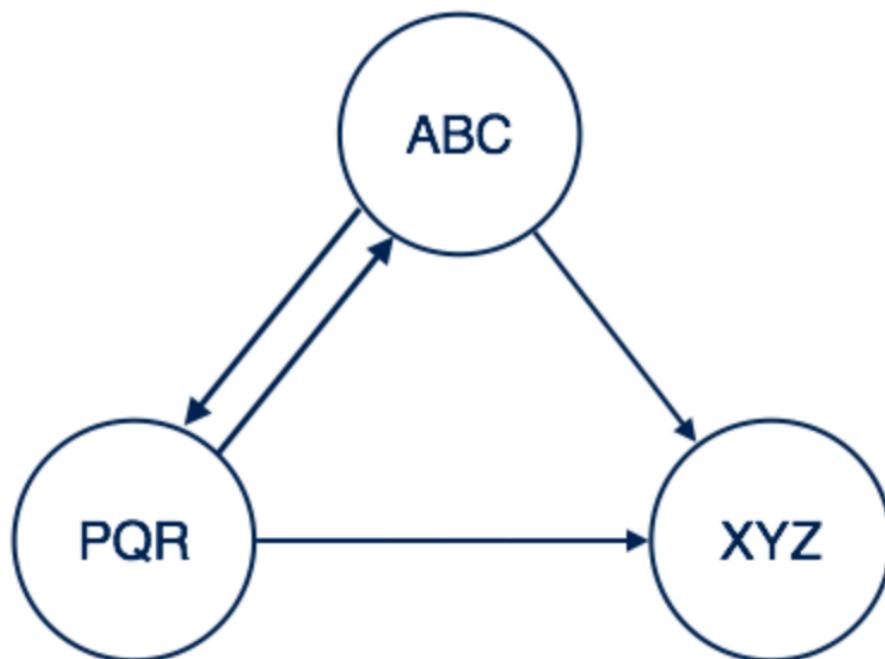
This node contains a set of properties. A Property is a name:value pair.

Relationships are created between two nodes



Here Relationship name "Follows" is created between two profiles. That means Profile-I follows Profile-II.

## Complex Graph Example



Here nodes are connected with relationships. Relationships are unidirectional or bi-directional.

- Relationship from XYZ to PQR is a unidirectional relationship.
- Relationship from ABC to XYZ is a bidirectional relationship.

The following are frequently used Neo4j CQL Commands/Clauses:

S.No.	CQL Command/Clause	Usage
1.	CREATE	To create nodes, relationships and properties
2.	MATCH	To retrieve data about nodes, relationships and properties
3.	RETURN	To return query results
4.	WHERE	To provide conditions to filter retrieval data
5.	DELETE	To delete nodes and relationships
6.	REMOVE	To delete properties of nodes and relationships
7.	ORDER BY	To sort retrieval data
8.	SET	To add or update labels

# Neo4j CQL Functions

The following are frequently used Neo4j CQL Functions:

S.No.	CQL Functions	Usage
1.	String	They are used to work with String literals.
2.	Aggregation	They are used to perform some aggregation operations on CQL Query results.
3.	Relationship	They are used to get details of Relationships like startnode,endnode etc.

CREATE command

CREATE (<node-name>:<label-name>)

Ejemplo: CREATE (emp:Employee)

-----

```
CREATE (
    <node-name>:<label-name>
    {
        <Property1-name>:<Property1-Value>
        .....
        <Propertyn-name>:<Propertyn-Value>
    }
)
```

CREATE (emp:Employee{ id:123, name:"Lokesh", sal:35000, deptno:10 })

CREATE (dept:Dept{ deptno:10, dname:"Accounting", location:"Hyderabad" })

Neo4j CQL MATCH command is used -

- To get data about nodes and properties from database
- To get data about nodes, relationships and properties from database

```
MATCH
(
    <node-name>:<label-name>
)
RETURN
    <node-name>.<property1-name>,
    .....
    <node-name>.<propertyn-name>
```

Ejemplo

```
MATCH (dept:Dept)
RETURN dept.deptno, dept.dname
```

```
MATCH (dept:Dept)
RETURN dept
```



## Relaciones entre nodos (relationships)

```
CREATE (e:Customer{id:"1001",name:"Abc",dob:"01/10/1982"})
```

```
CREATE
(cc:CreditCard{id:"5001",number:"1234567890",cvv:"888",expireddate:"20/17"})
----
```

```
MATCH (e:Customer), (cc:CreditCard)
CREATE (e) -[r:DO_SHOPPING_WITH] -> (cc)
```

```
MATCH (cust:Customer), (cc:CreditCard)
CREATE (cust) -[r:DO_SHOPPING_WITH{shopdate:"12/12/2014", price:55000}] -> (cc)
RETURN r
```

```
MATCH (cust:Customer), (cc:CreditCard)
WHERE cust.id = "1001" AND cc.id= "5001"
CREATE (cust) -[r:DO_SHOPPING_WITH{shopdate:"12/12/2014", price:55000}] -> (cc)
RETURN r
```

```
CREATE
```

```
(video1:YoutubeVideo1 {  
    title:"Action Movie1",  
    updated_by:"Abc",  
    uploaded_date:"10/10/2010"  
}  
) -[movie:ACTION_MOVIES{rating:1}]->  
(video2:YoutubeVideo2 {  
    title:"Action Movie2",  
    updated_by:"Xyz",  
    uploaded_date:"12/12/2012"  
}  
)
```

```
-----
```

```
MATCH (emp:Employee)  
WHERE emp.name = 'Abc'  
RETURN emp
```

```
MATCH (emp:Employee)  
WHERE emp.name = 'Abc' OR emp.name = 'Xyz'  
RETURN emp
```

Neo4j CQL REMOVE command is used

- To remove labels of a Node or a Relationship
- To remove properties of a Node or a Relationship

Main Difference between Neo4j CQL DELETE and REMOVE commands -

- DELETE operation is used to delete Nodes and associated Relationships.
- REMOVE operation is used to remove labels and properties.

```
CREATE (book:Book {id:122,title:"Neo4j Tutorial",pages:340,price:250})
```

```
MATCH (book { id:122 })
REMOVE book.price
RETURN book
```

Equivalent SQL:

```
ALTER TABLE BOOK REMOVE COLUMN PRICE;
SELECT * FROM BOOK WHERE ID = 122;
```

Agrega un nuevo campo (par propiedad-valor), si ya existe lo actualiza al nuevo valor.

```
MATCH (dc:DebitCard)
SET dc.atm_pin = 3456
RETURN dc
```

-----

```
MATCH (emp:Employee)
RETURN emp.empid, emp.name, emp.salary, emp.deptno
ORDER BY emp.name
```

-----

```
MATCH (emp:Employee)
RETURN emp.empid, emp.name, emp.salary, emp.deptno
ORDER BY emp.name DESC
```

-----

```
MATCH (cc:CreditCard) RETURN cc.id, cc.number
UNION
MATCH (dc:DebitCard) RETURN dc.id, dc.number
```

```
MATCH (cc:CreditCard)
RETURN cc.id as id,cc.number as number,cc.name as name,
      cc.valid_from as valid_from,cc.valid_to as valid_to
UNION
MATCH (dc:DebitCard)
RETURN dc.id as id,dc.number as number,dc.name as name,
      dc.valid_from as valid_from,dc.valid_to as valid_to
```

-----

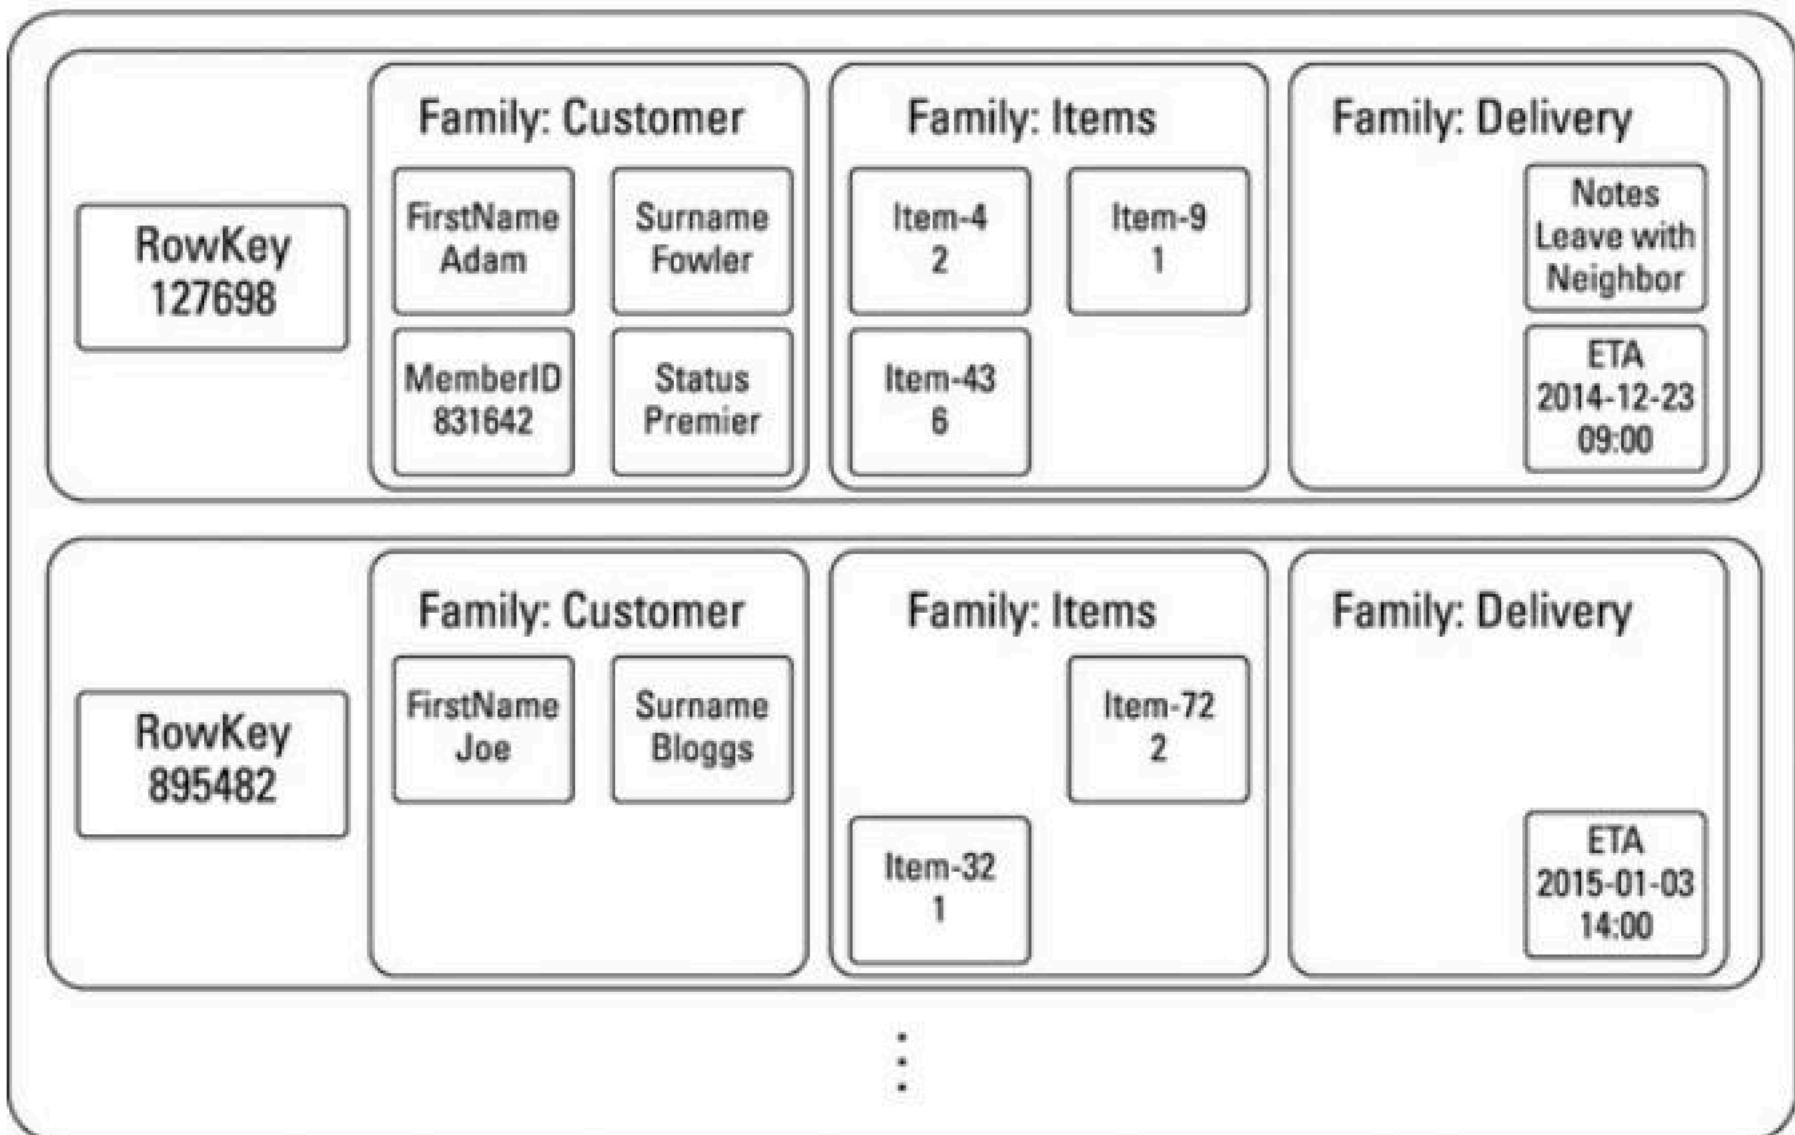
```
MATCH (e:Employee)
WHERE e.id IS NOT NULL
RETURN e.id,e.name,e.sal,e.deptno
```

-----

```
MATCH (e:Employee)
WHERE e.id IN [123,124]
RETURN e.id,e.name,e.sal,e.deptno
```

# Bases de datos Columnares

Order Table



# Query Driven Methodology: process

- Entities and relationships: map to tables
- Key attributes: map to primary key columns
- Equality search attributes: must be at the beginning of the primary key
- Inequality search attributes: become clustering columns
- Ordering attributes: become clustering columns

PARTITION  
KEY +

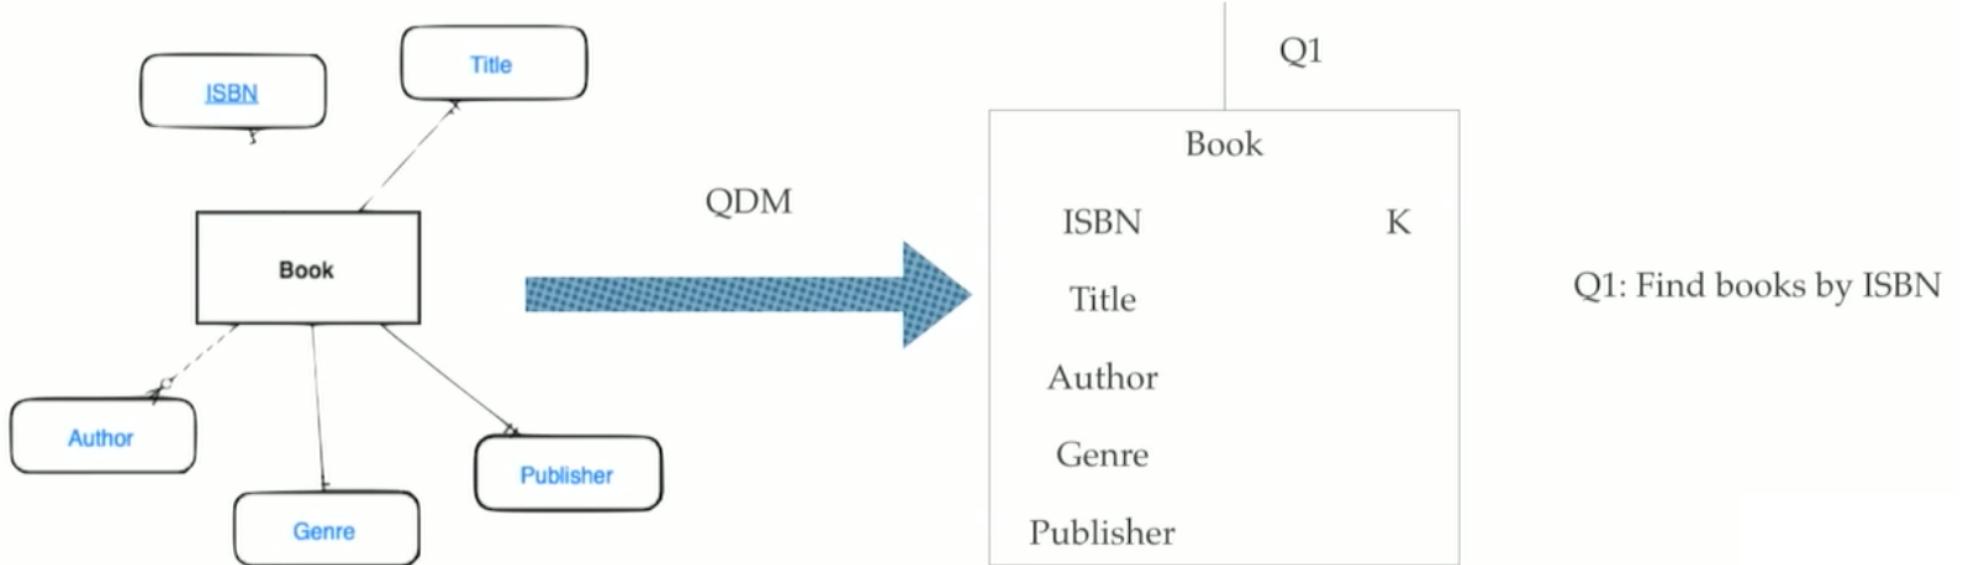
CLUSTERING  
COLUMN(S)

(dos elementos para obtener rendimiento eficiente)

El diseño del modelo de datos está condicionado por las consultas que realiza la aplicación en la base de datos (Column Families).

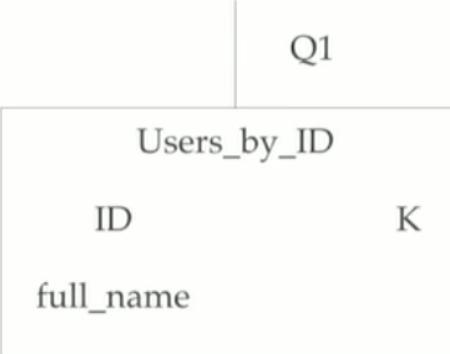
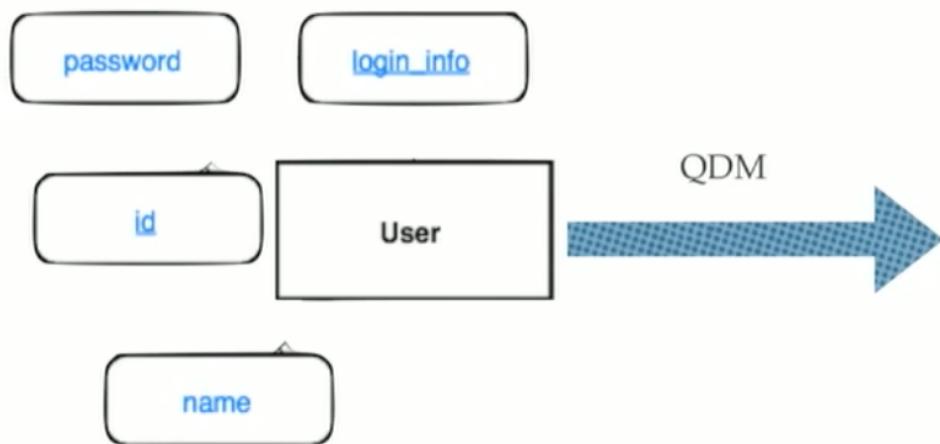
# Requirement: 1

Books can be uniquely identified and accessed by ISBN, we also need a title, genre, author and publisher.

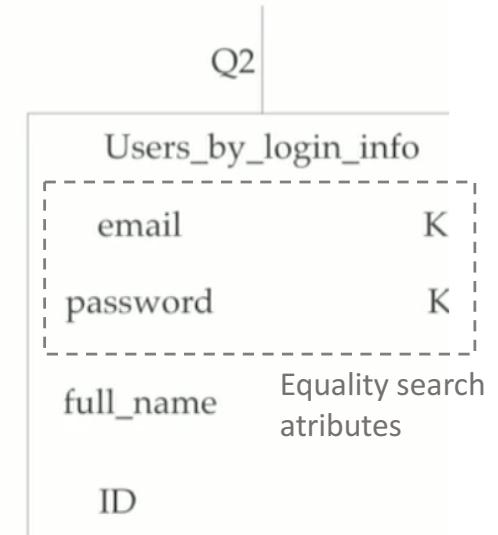


# Requirement 2

Users register into the system uniquely identified by an email and a password. We also want their full name. They will be accessed by email and password or internal unique ID.



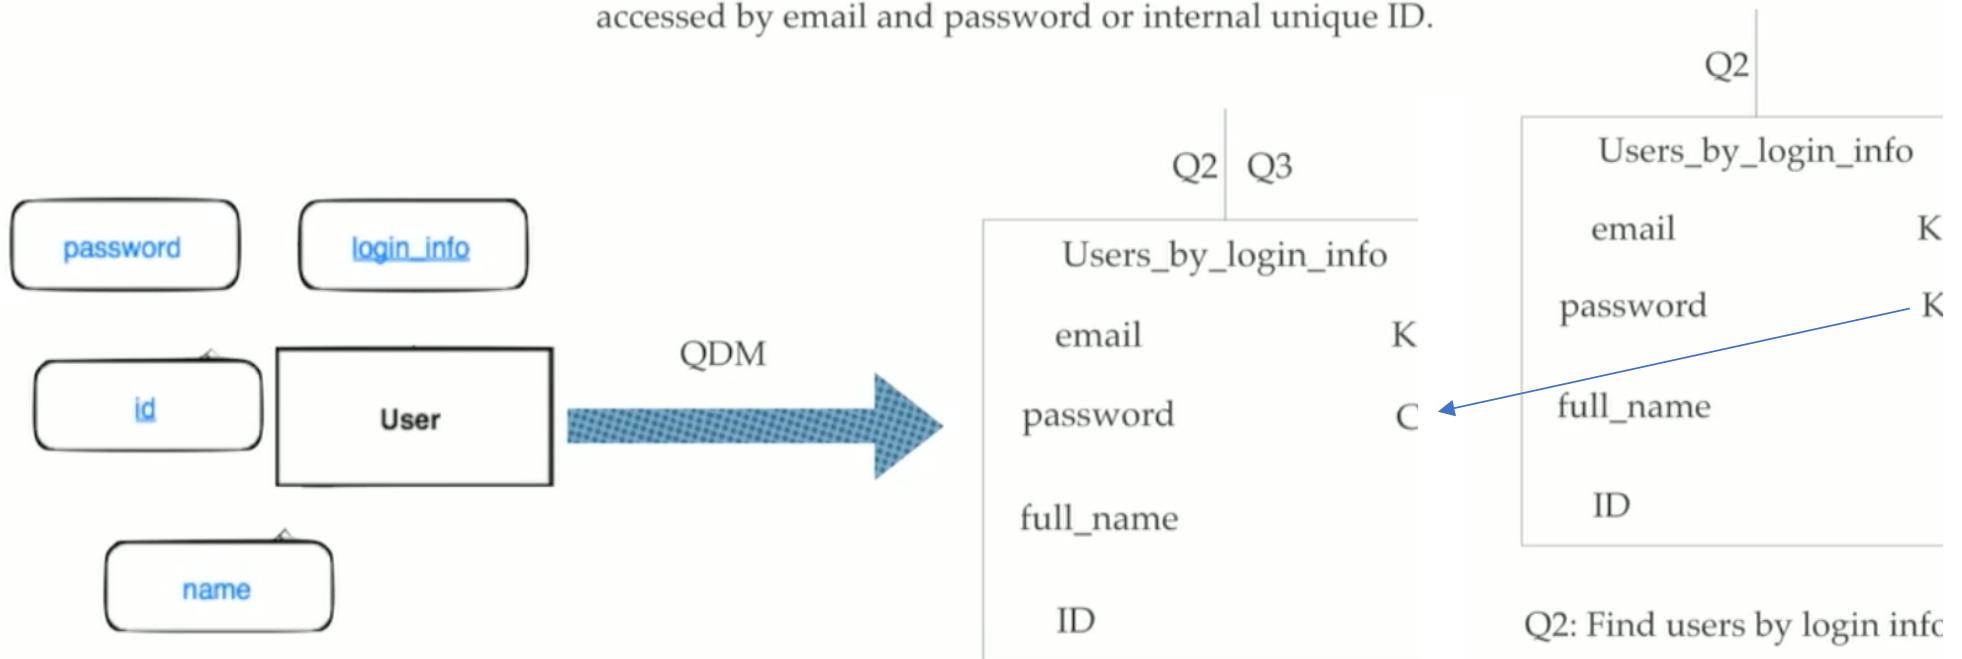
Q1: Find users by ID



Q2: Find users by login info

# Requirement 2

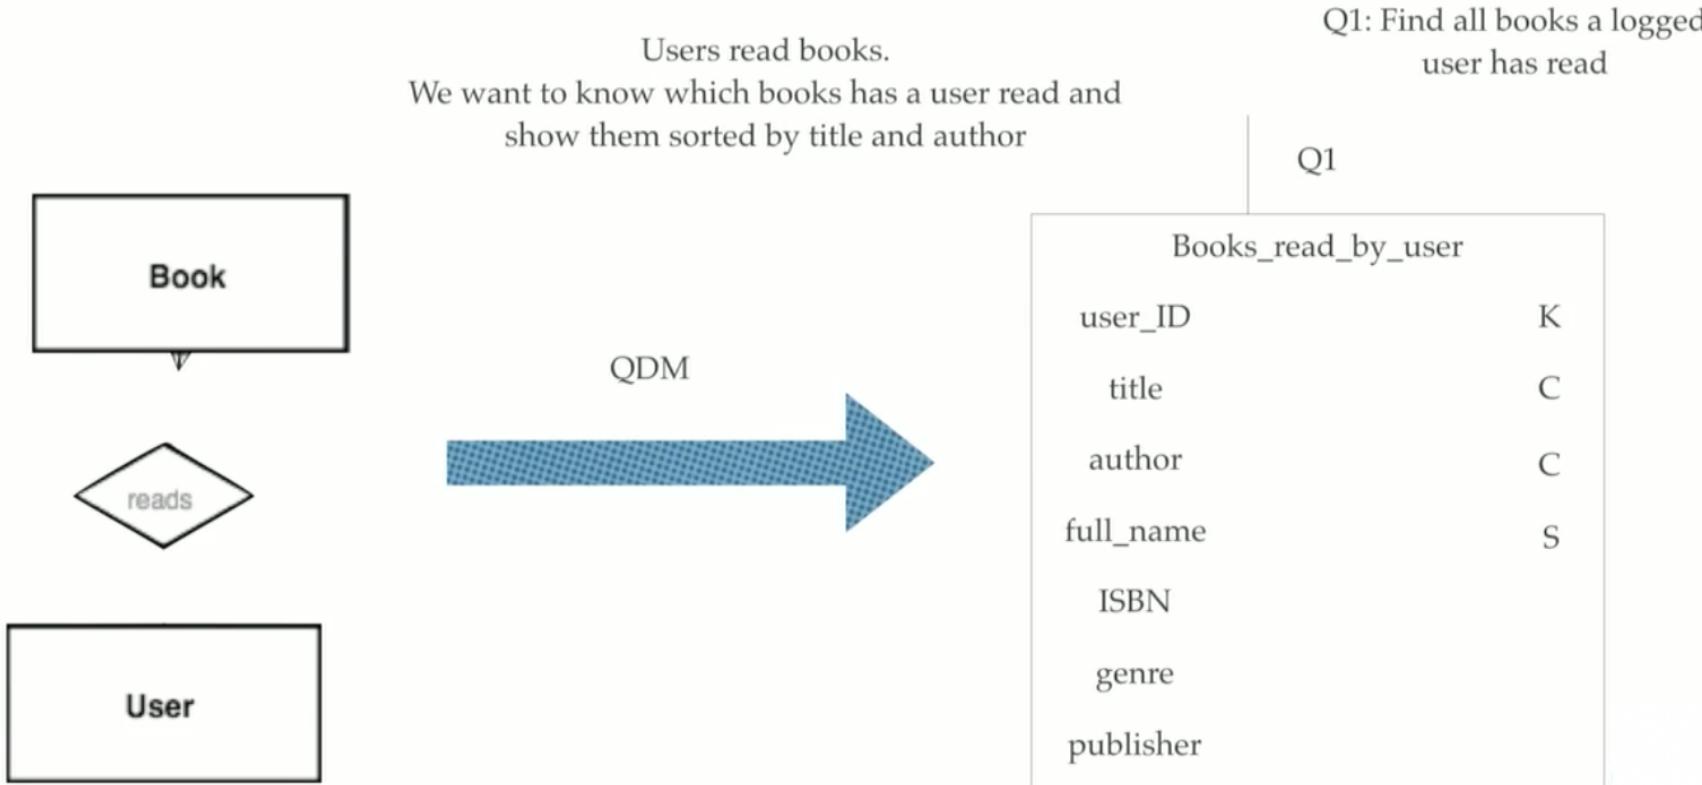
Users register into the system uniquely identified by an email and a password. We also want their full name. They will be accessed by email and password or internal unique ID.



La misma tabla se puede utilizar para acomodar consultas de nuevos usuarios (Q3).

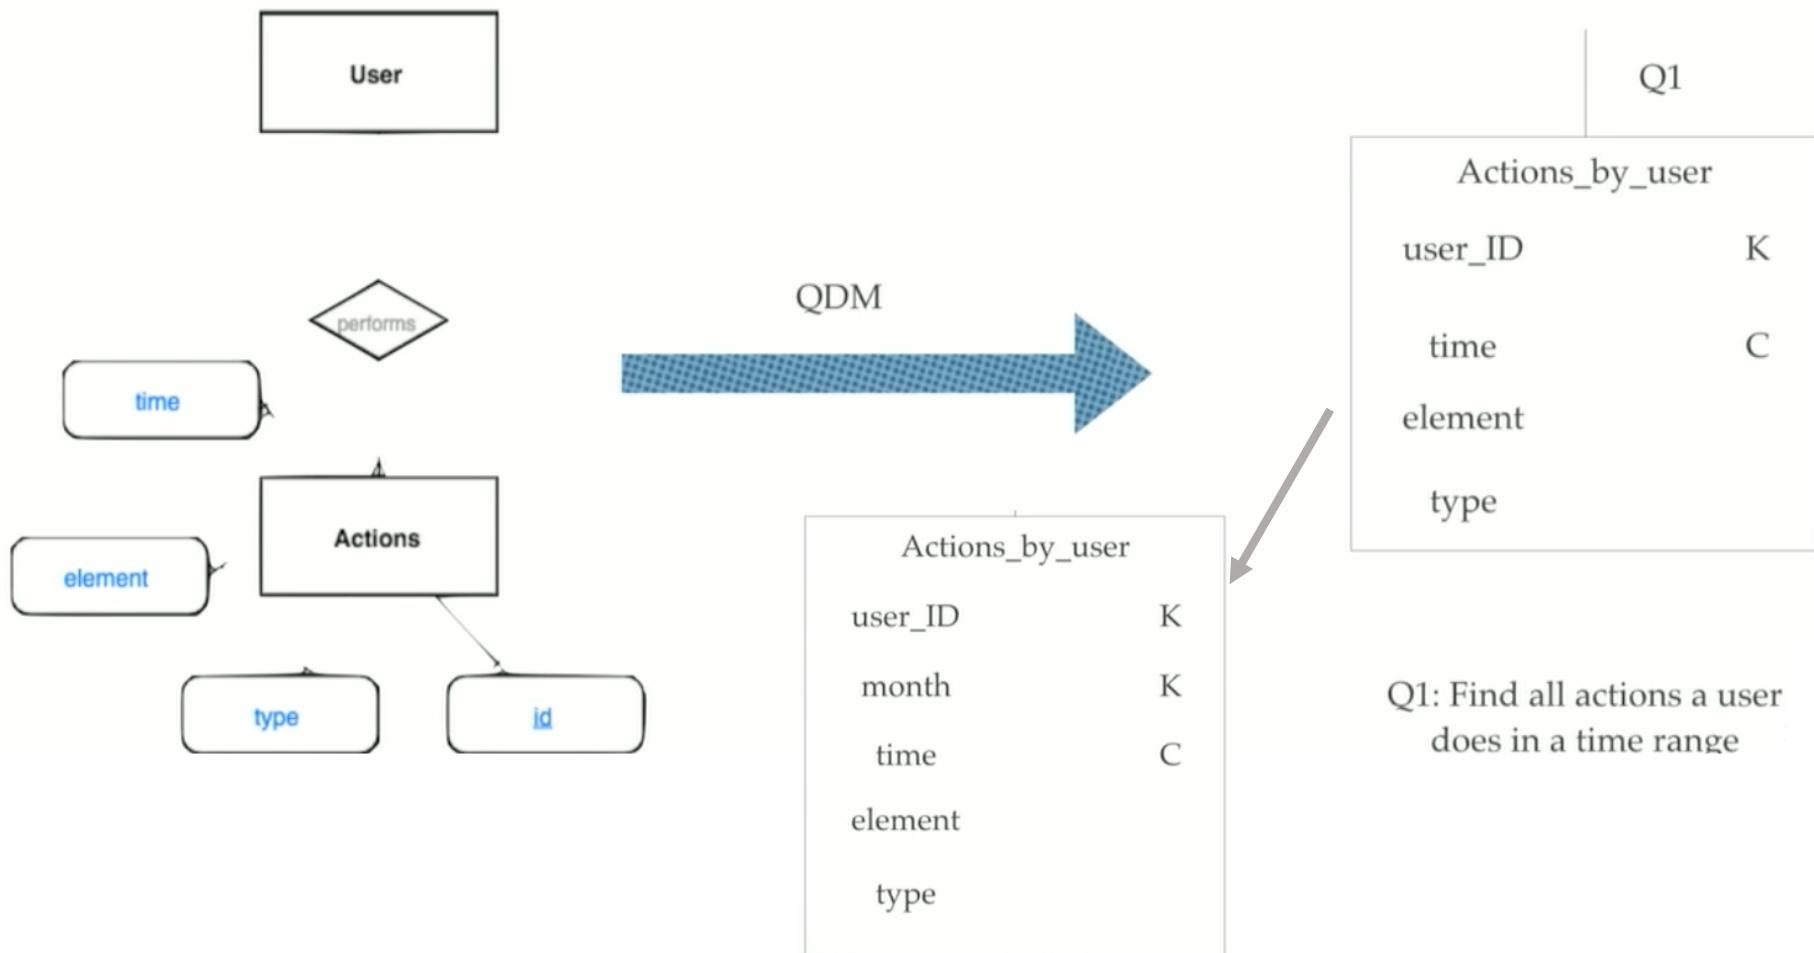
```
BEGIN BATCH
    INSERT INTO users_by_id (ID, full_name) VALUES (...) IF NOT EXISTS;
    INSERT INTO users_by_login_info (email, password, full_name, ID) VALUES (...);
APPLY BATCH;
```

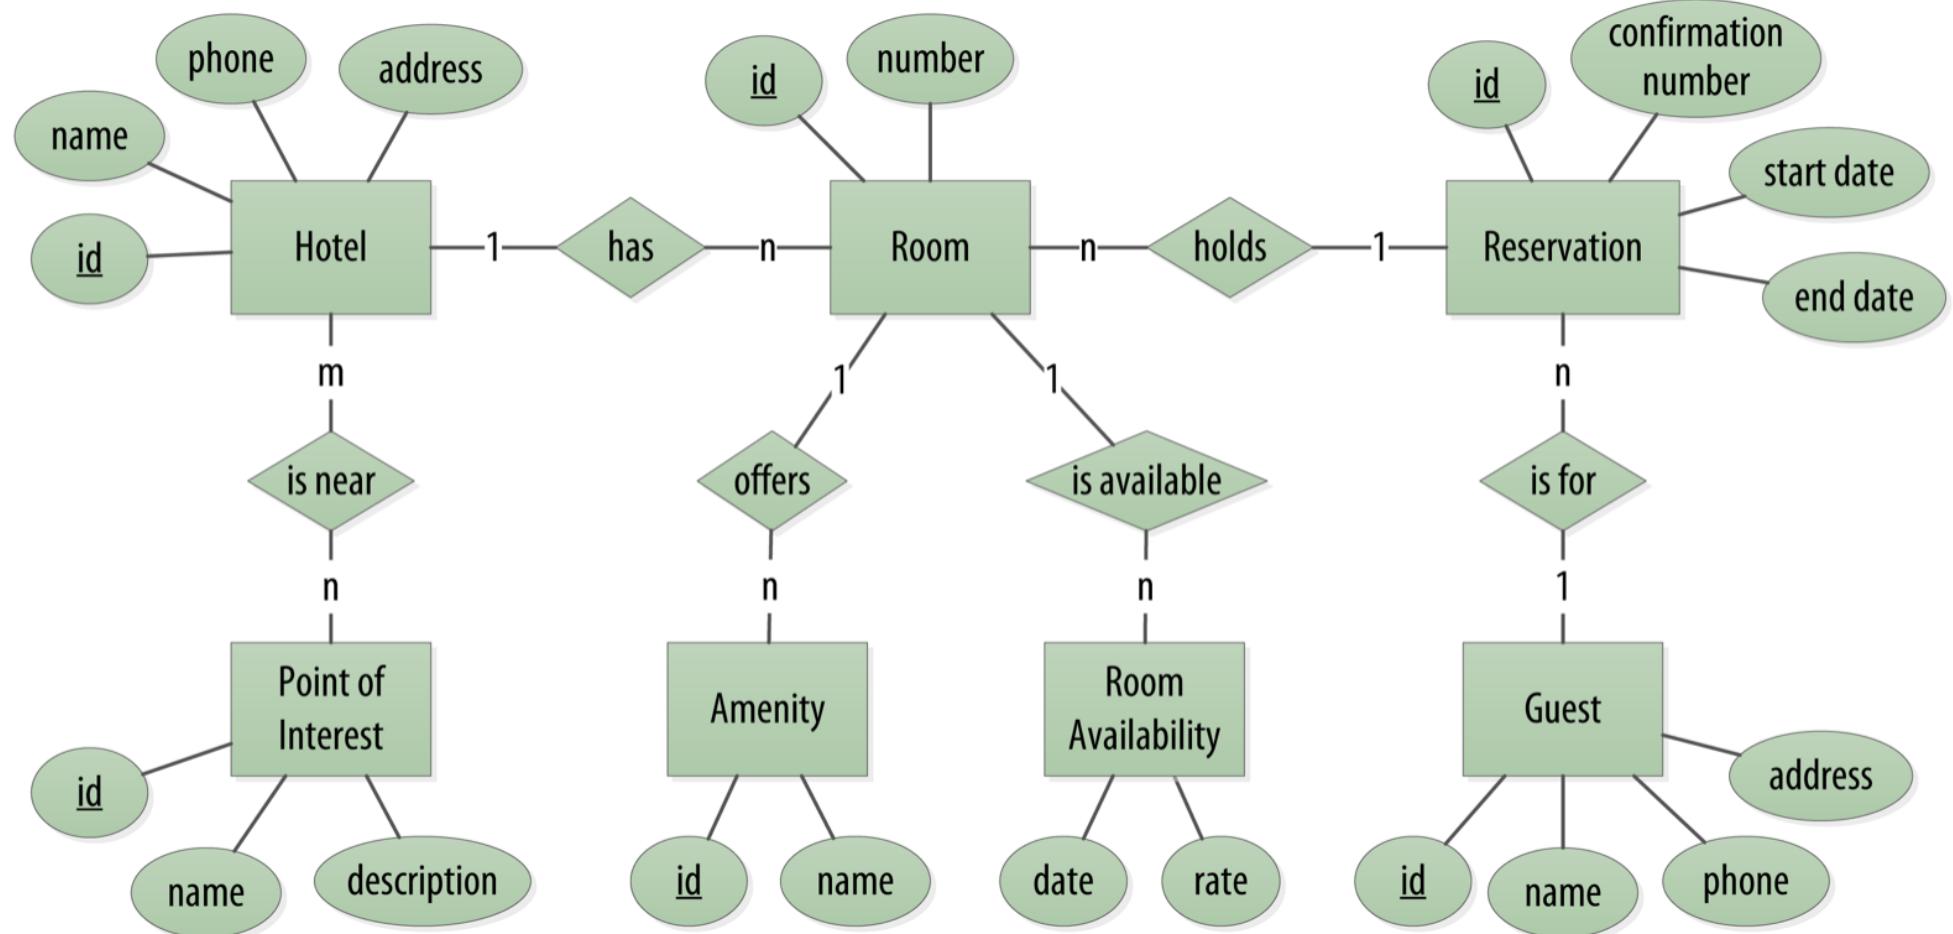
# Requirement 3

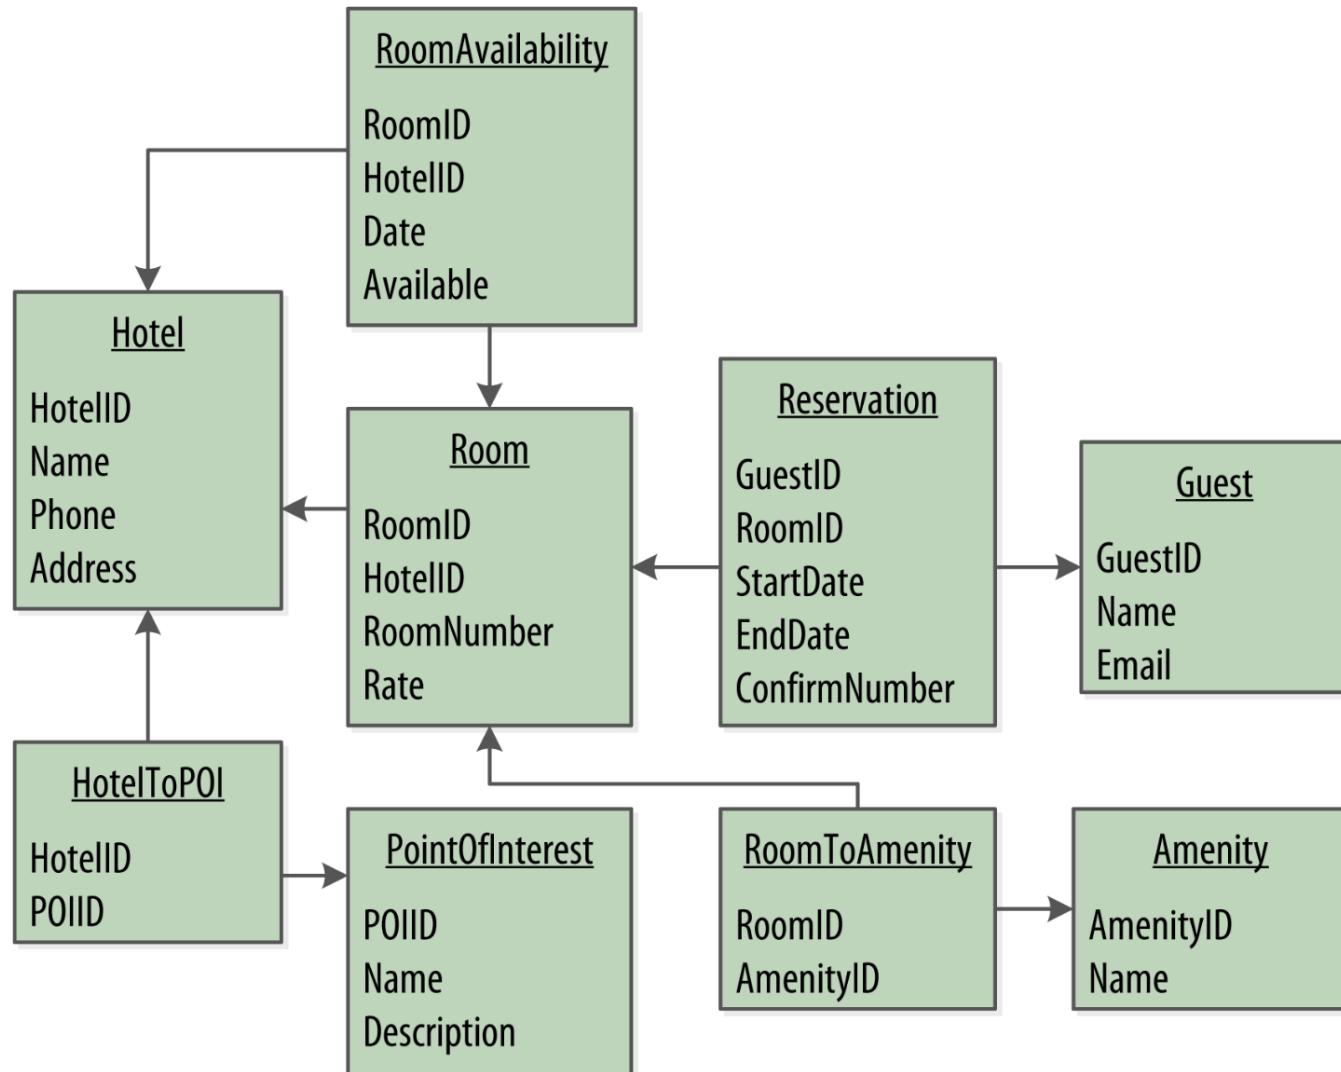


# Requirement 4

In order to improve our site's usability we need to understand how our users use it by tracking every interaction they have with our site.

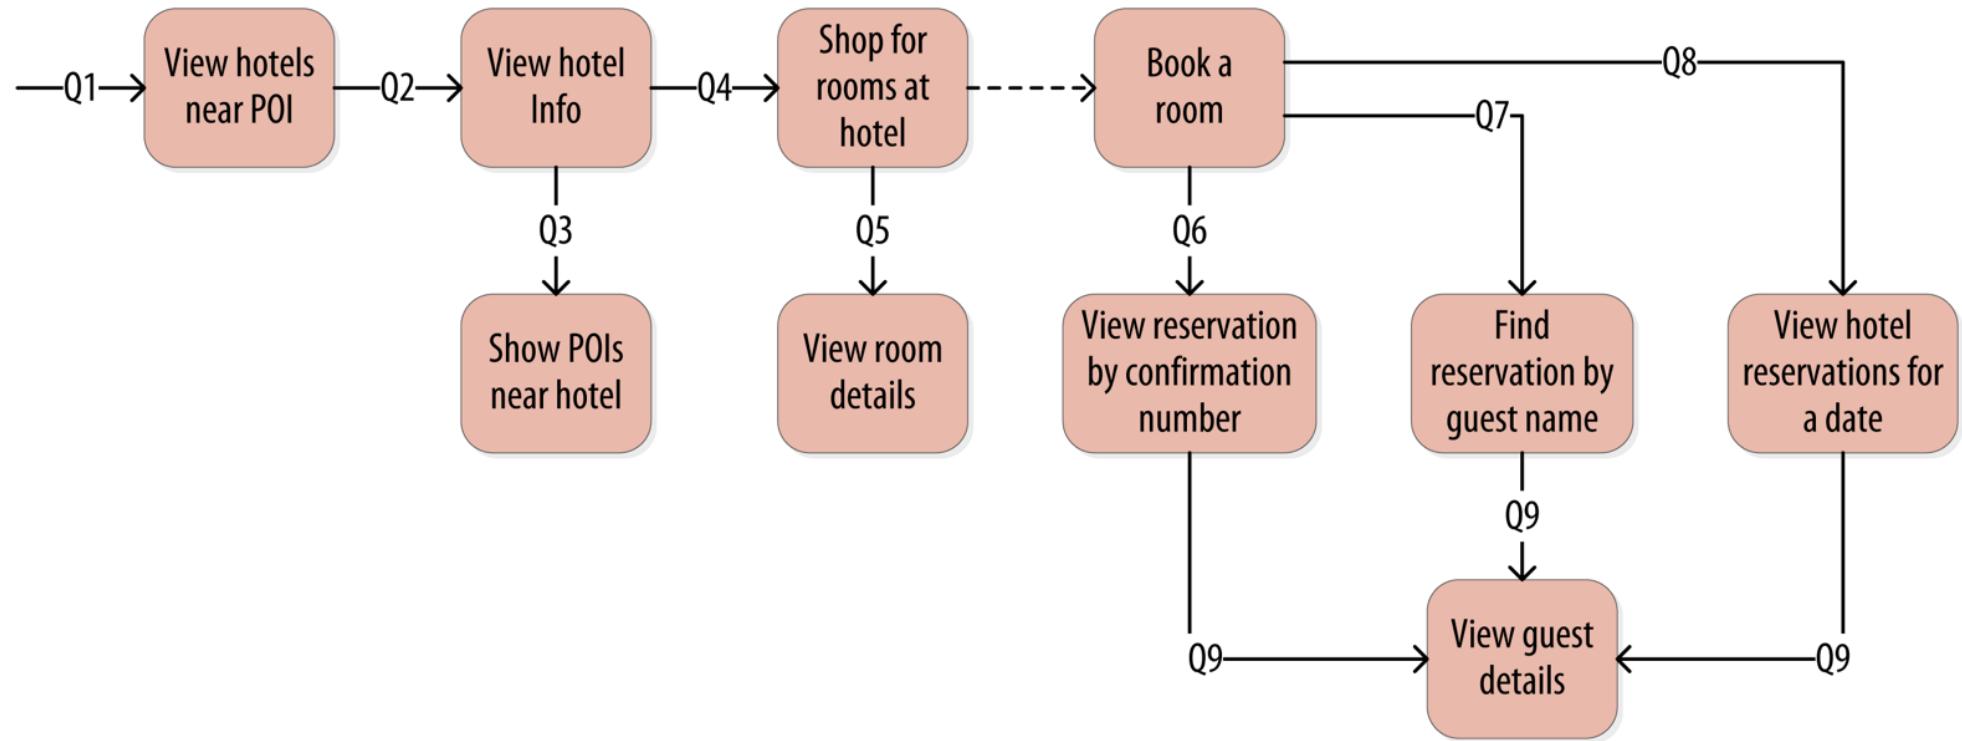




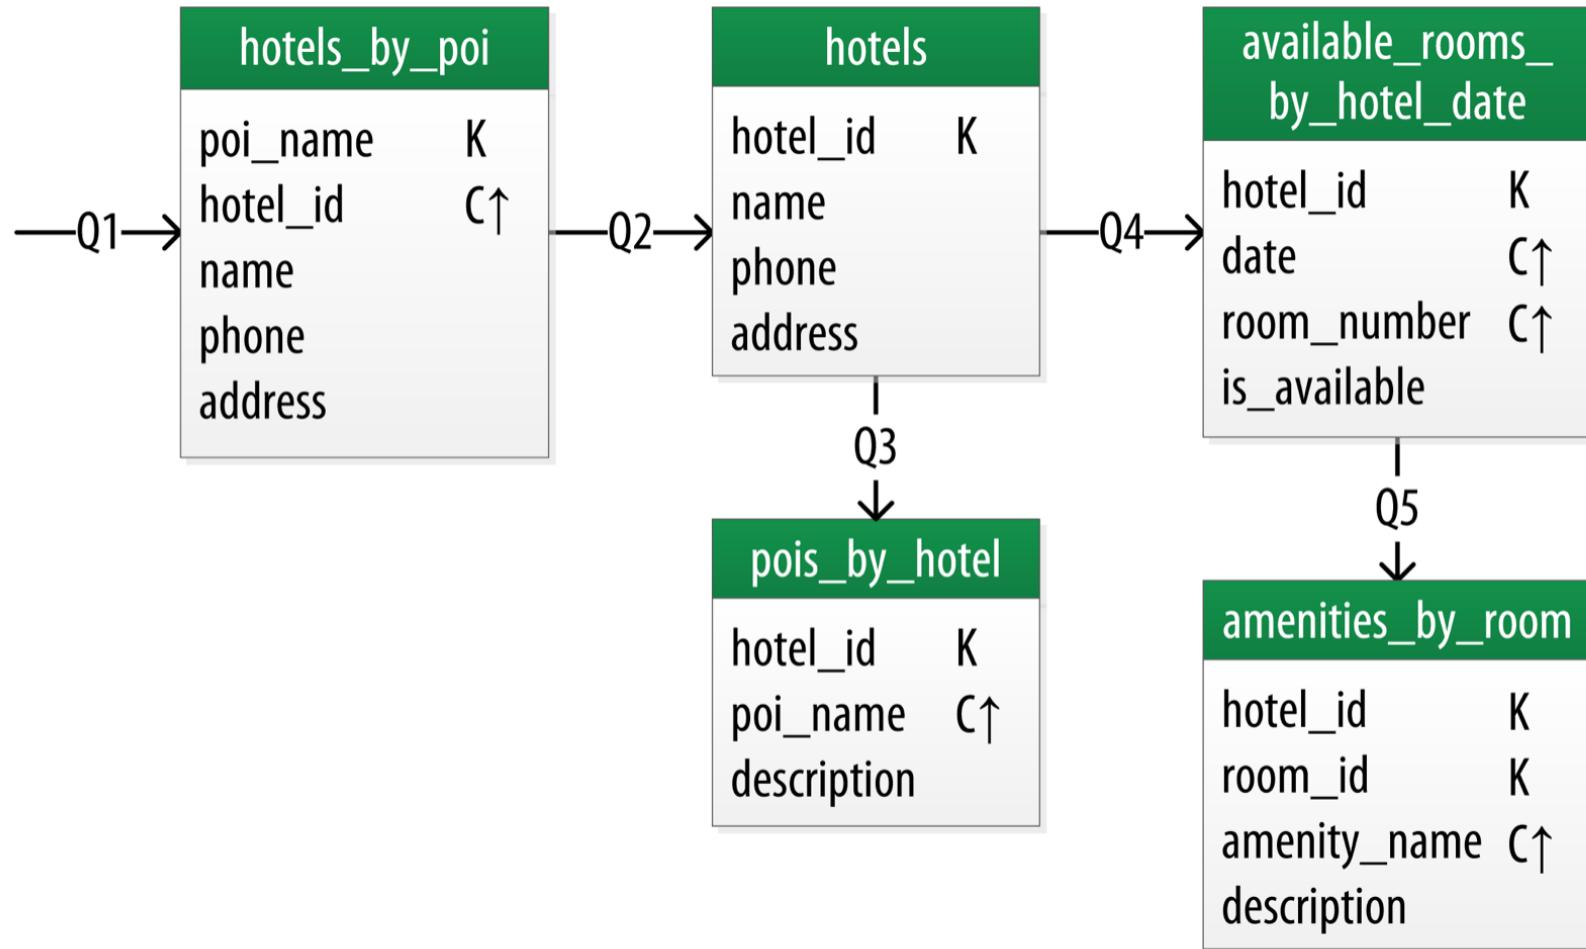


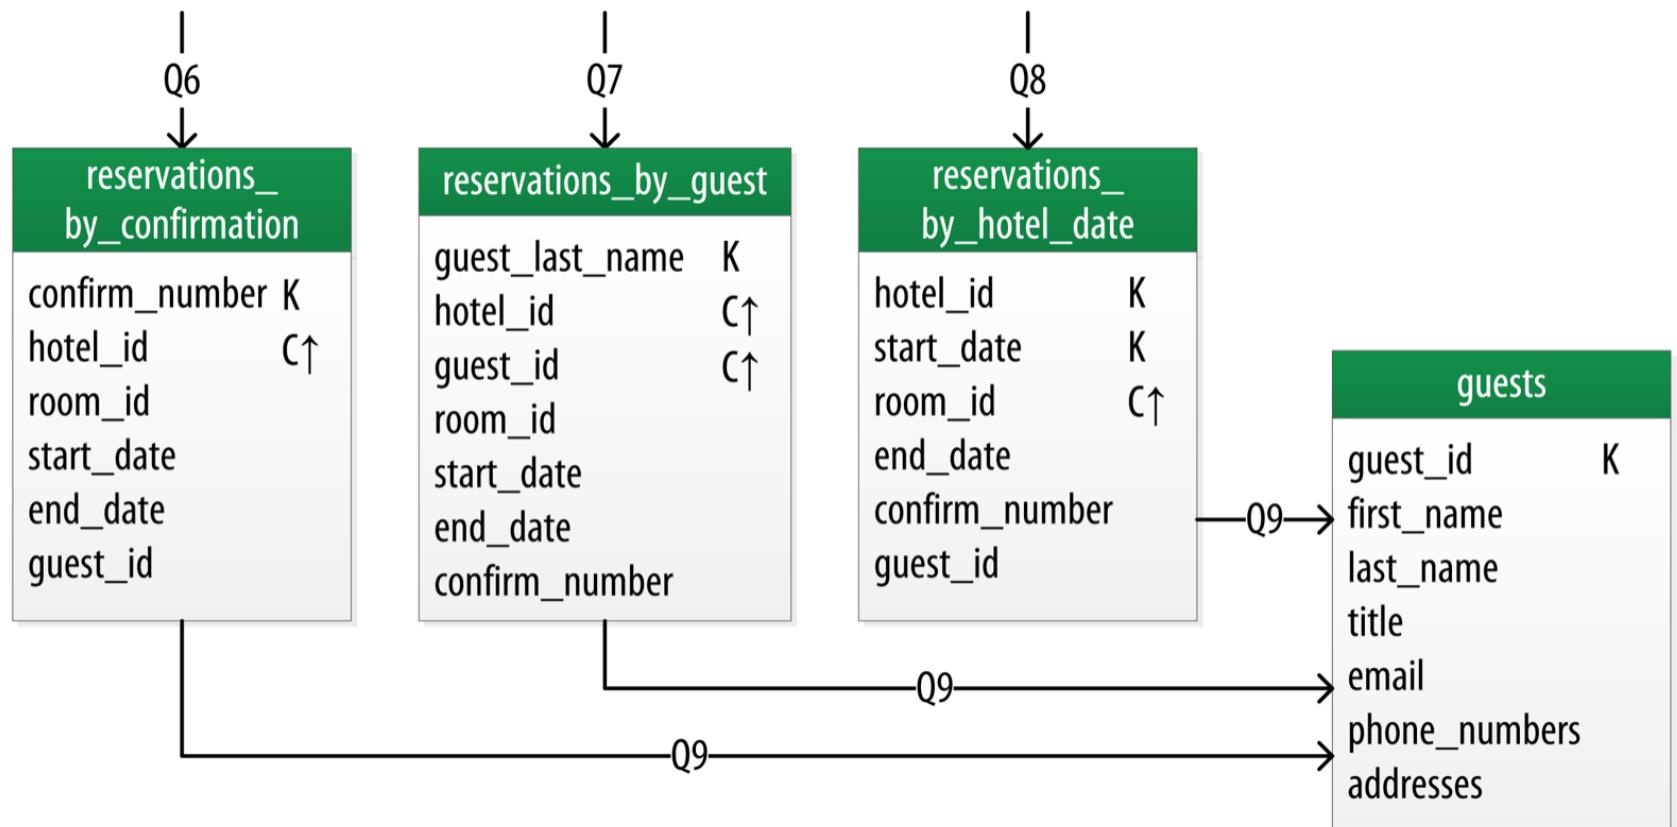
Modelo relacional de datos para el sistema de reserva de hoteles

- Q1. Find hotels near a given point of interest.
- Q2. Find information about a given hotel, such as its name and location.
- Q3. Find points of interest near a given hotel.
- Q4. Find an available room in a given date range.
- Q5. Find the rate and amenities for a room.
- Q6. Lookup a reservation by confirmation number.
- Q7. Lookup a reservation by hotel, date, and guest name.
- Q8. Lookup all reservations by guest name.
- Q9. View guest details.



Organización lógica de las consultas que debe proporcionar el sistema de reserva de hoteles.





## hotel keyspace

hotels		
hotel_id	text	K
name	text	
phone	text	
*address*	address	

hotels_by_poi		
poi_name	text	K
hotel_id	text	C↑
name	text	
phone	text	
address	address	

*address*		
street	text	
city	text	
state_or_province	text	
postal_code	text	
country	text	

available_rooms_by_hotel_date		
hotel_id	text	K
date	date	C↑
room_number	smallint	C↑
is_available	boolean	

pois_by_hotel		
hotel_id	text	K
poi_name	text	C↑
description	text	

amenities_by_room		
hotel_id	text	K
room_number	smallint	K
amenity_name	text	C↑
description	text	

```
CREATE KEYSPACE hotel
    WITH replication = {'class': 'SimpleStrategy', 'replication_factor' : 3};

CREATE TYPE hotel.address (
    street text,
    city text,
    state_or_province text,
    postal_code text,
    country text
);

CREATE TABLE hotel.hotels_by_poi (
    poi_name text,
    hotel_id text,
    name text,
    phone text,
    address frozen<address>,
    PRIMARY KEY ((poi_name), hotel_id)
) WITH comment = 'Q1. Find hotels near given poi'
AND CLUSTERING ORDER BY (hotel_id ASC) ;
```

## reservation keyspace

reservations_by_hotel_date		
hotel_id	text	K
start_date	date	K
room_number	smallint	C↑
nights	smallint	
confirm_number	text	
guest_id	uuid	

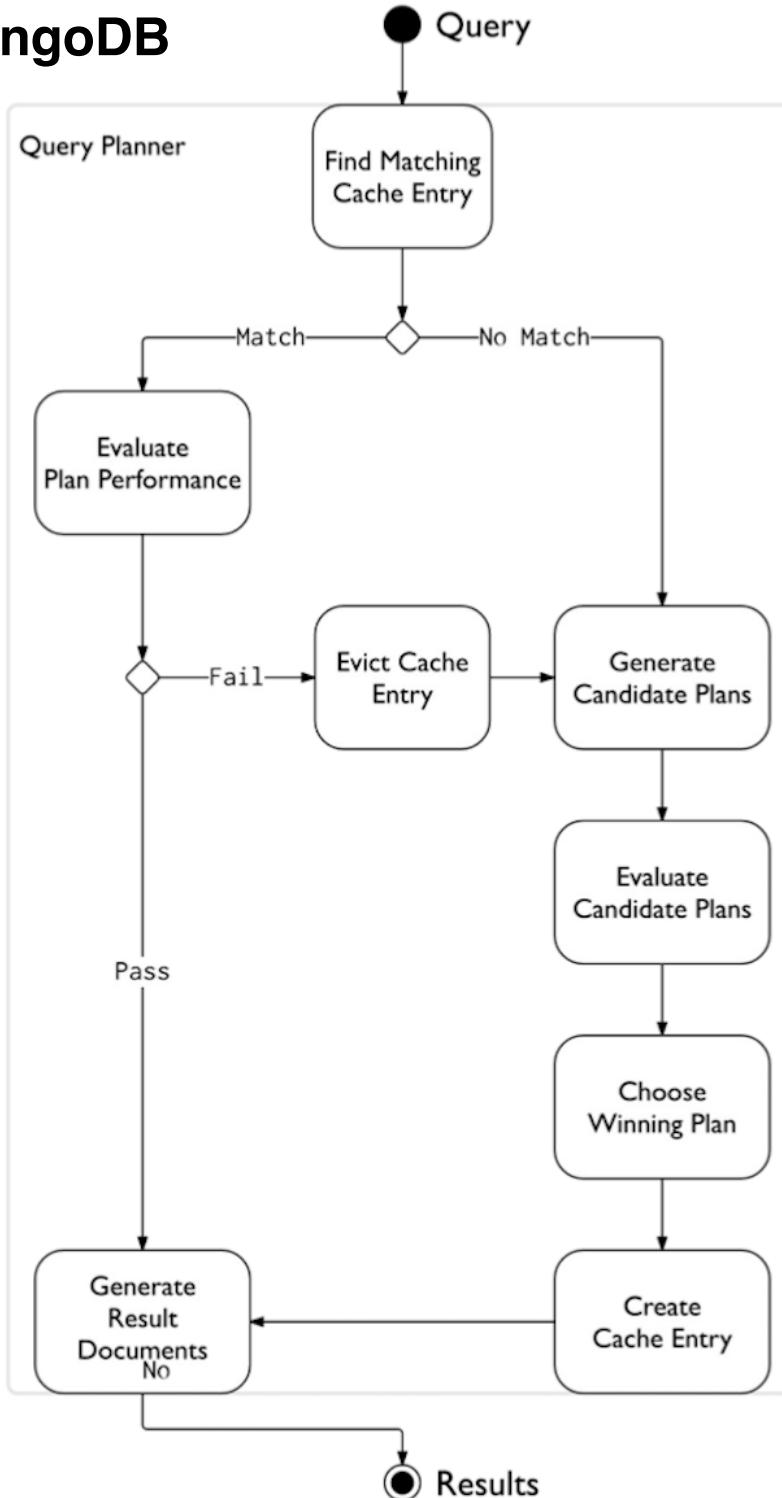
reservations_by_confirmation		
confirm_number	text	K
hotel_id	text	C↑
start_date	date	C↑
room_number	smallint	C↑
nights	smallint	
guest_id	uuid	

reservations_by_guest		
guest_last_name	text	K
hotel_id	text	C↑
room_number	smallint	C↑
start_date	date	C↑
nights	smallint	
confirm_number	text	
guest_id	uuid	

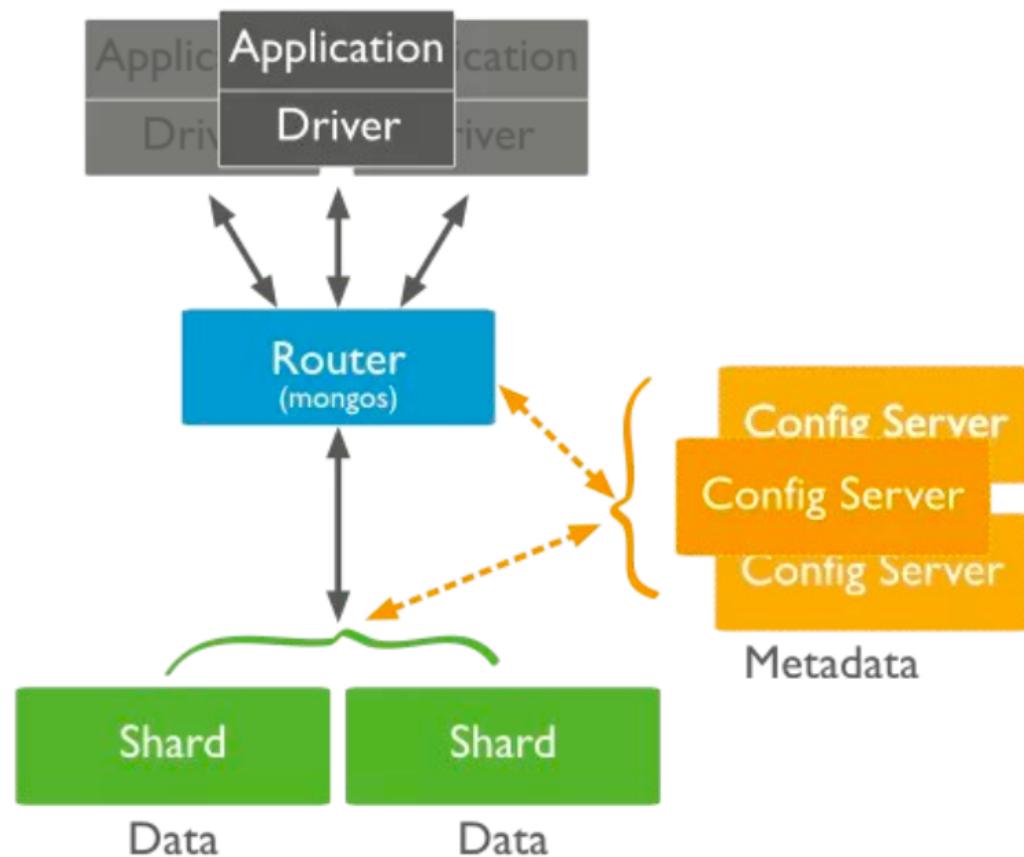
guests		
guest_id	uuid	K
first_name	text	
last_name	text	
title	text	
{emails}	text	
[phone_numbers]	text	
<addresses>	text, address	

*address*	
street	text
city	text
state_or_province	text
postal_code	text
country	text

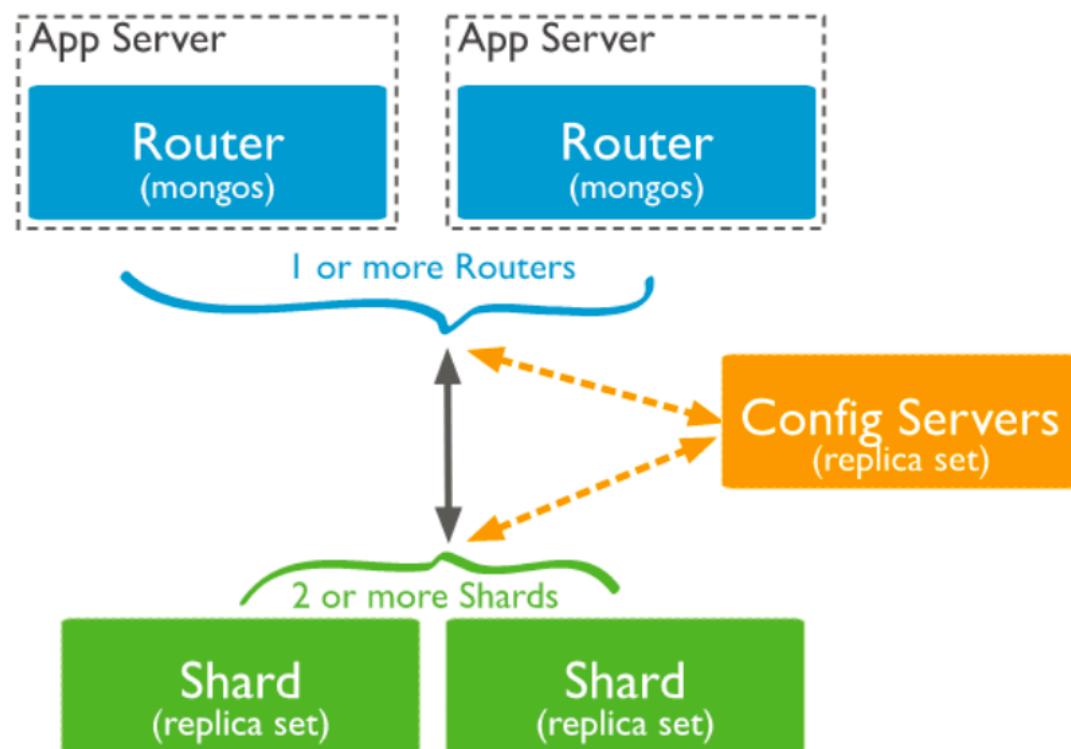
# Escalabilidad de MongoDB

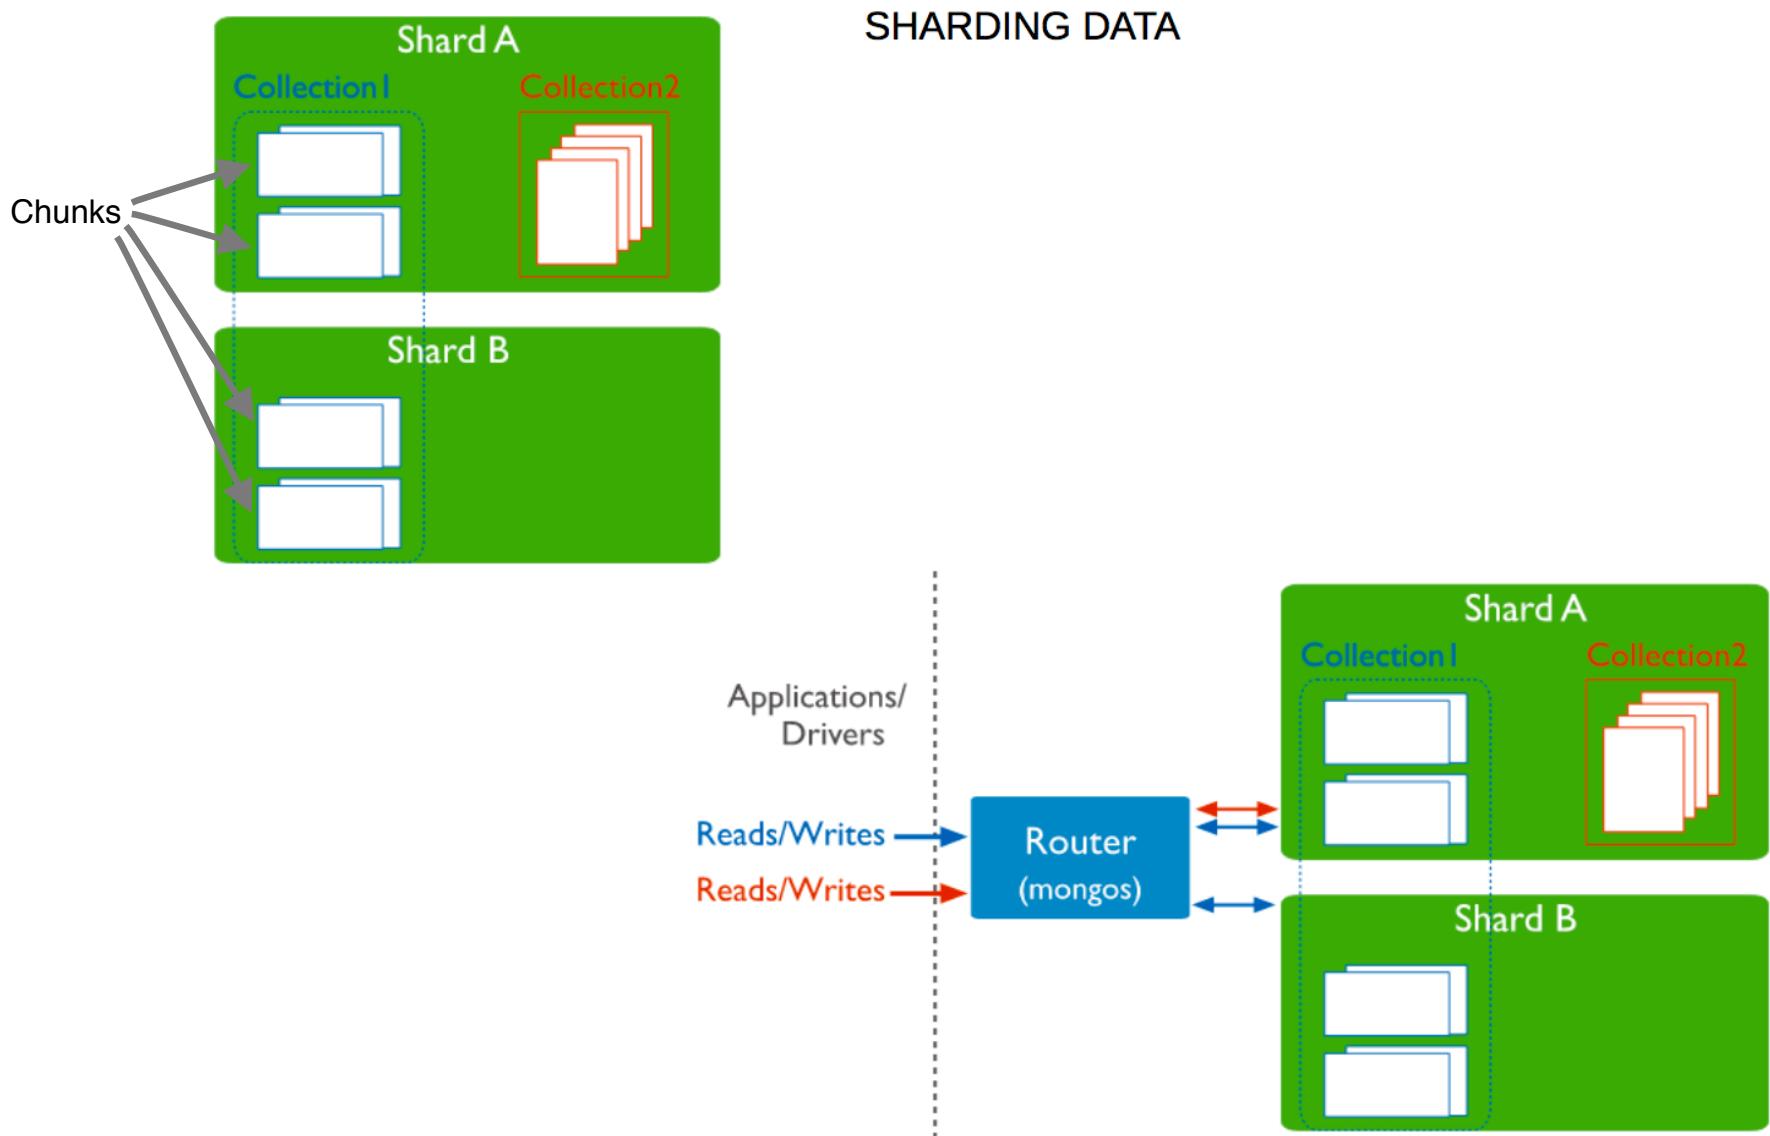


## ARCHITECTURE MONGODB



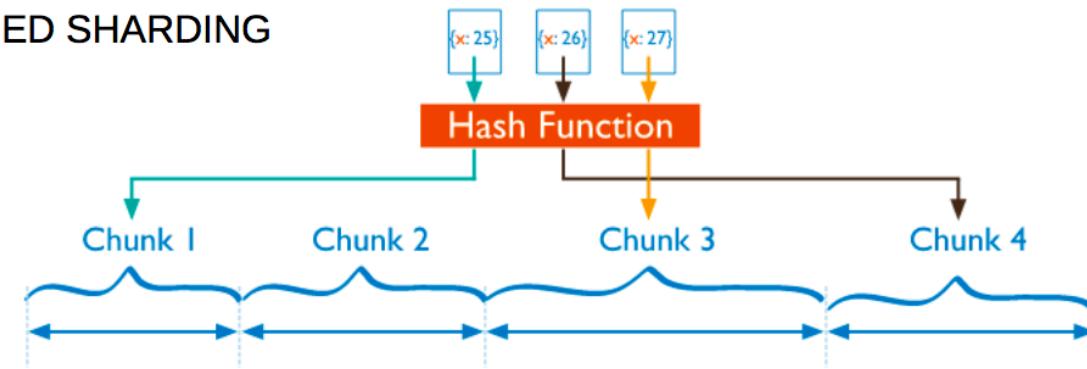
## SHARD CLUSTER (SHARDING DATA)



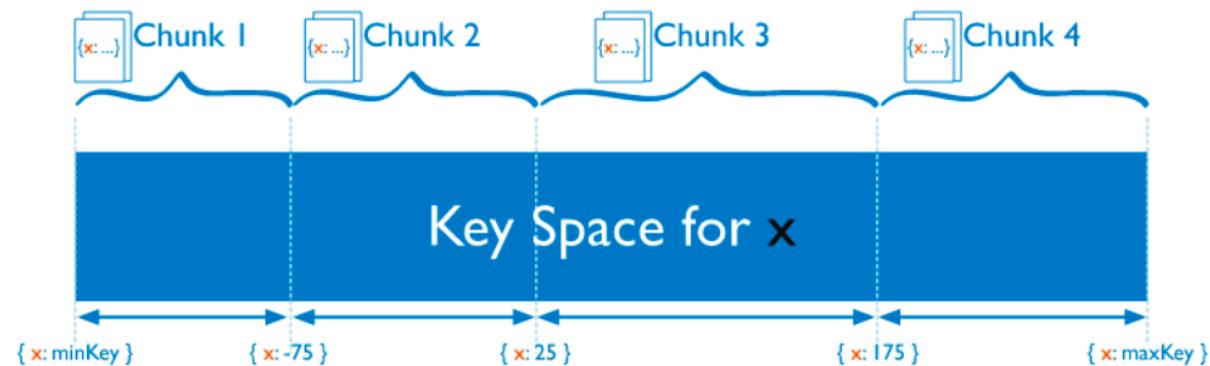


## DISTRIBUTED SHARDING

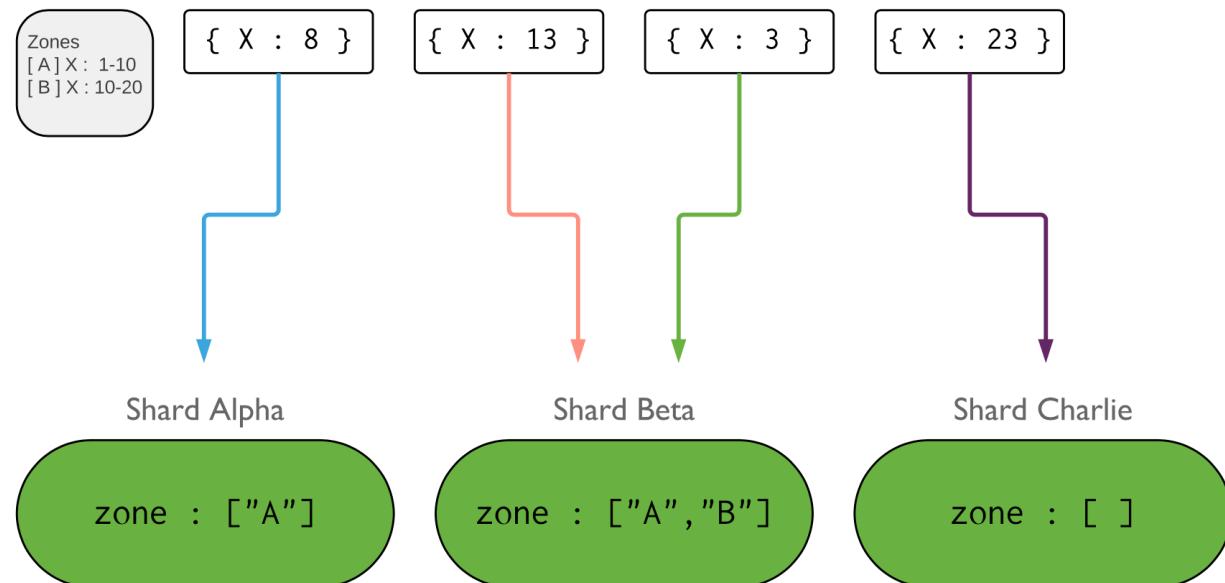
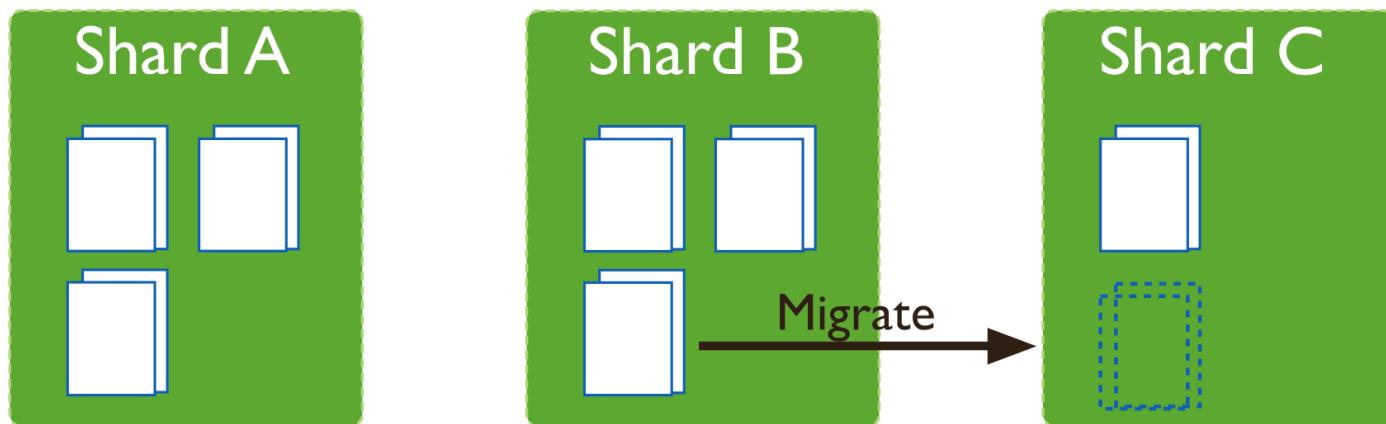
### HASHED SHARDING



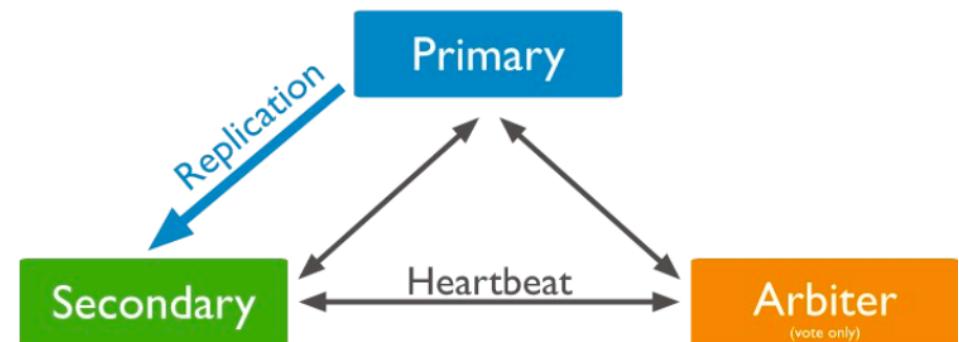
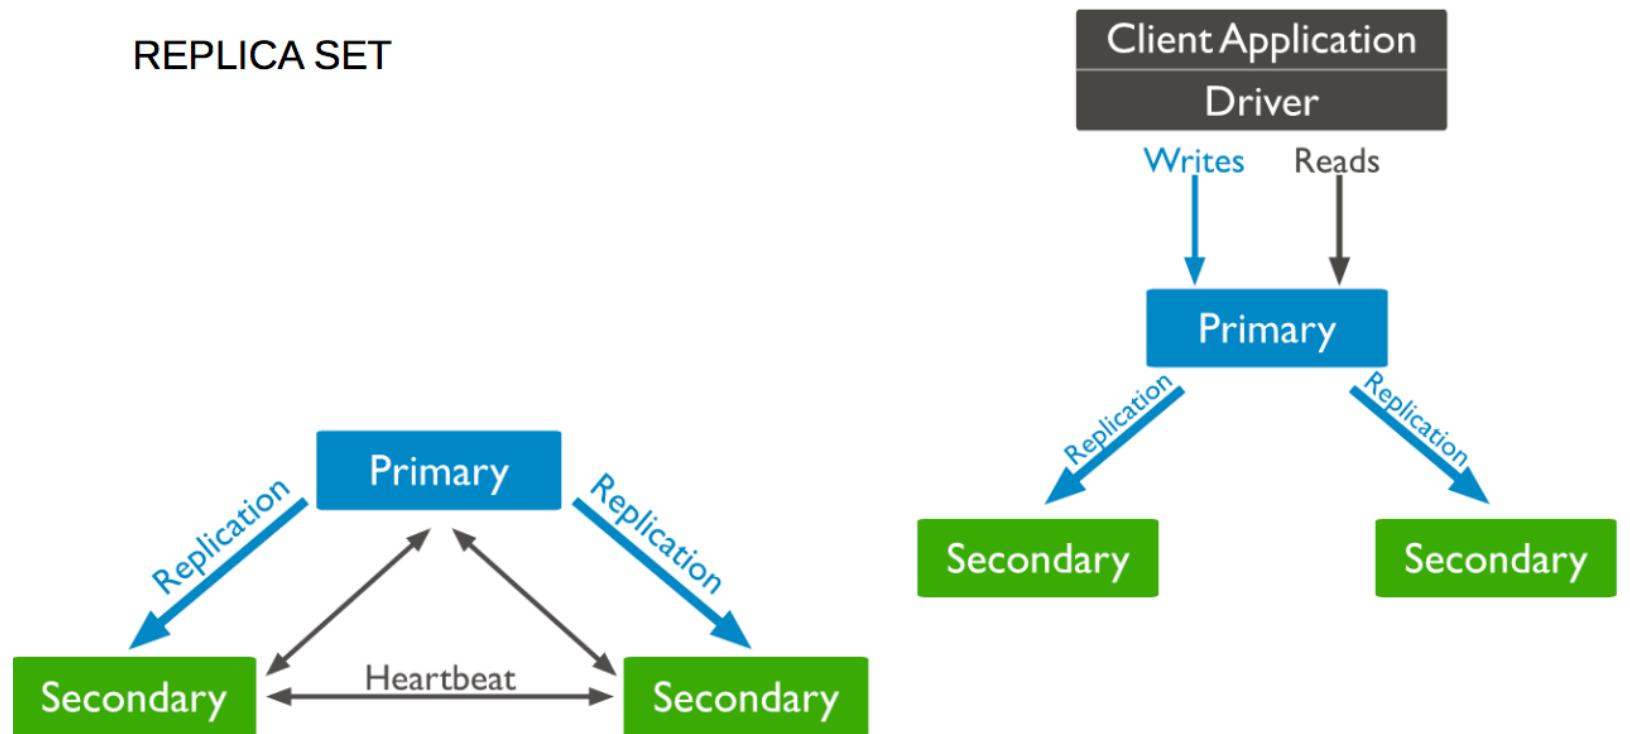
### RANGED SHARDING



El balanceador de carga opera re-distribuyendo los "chunks"



## REPLICA SET



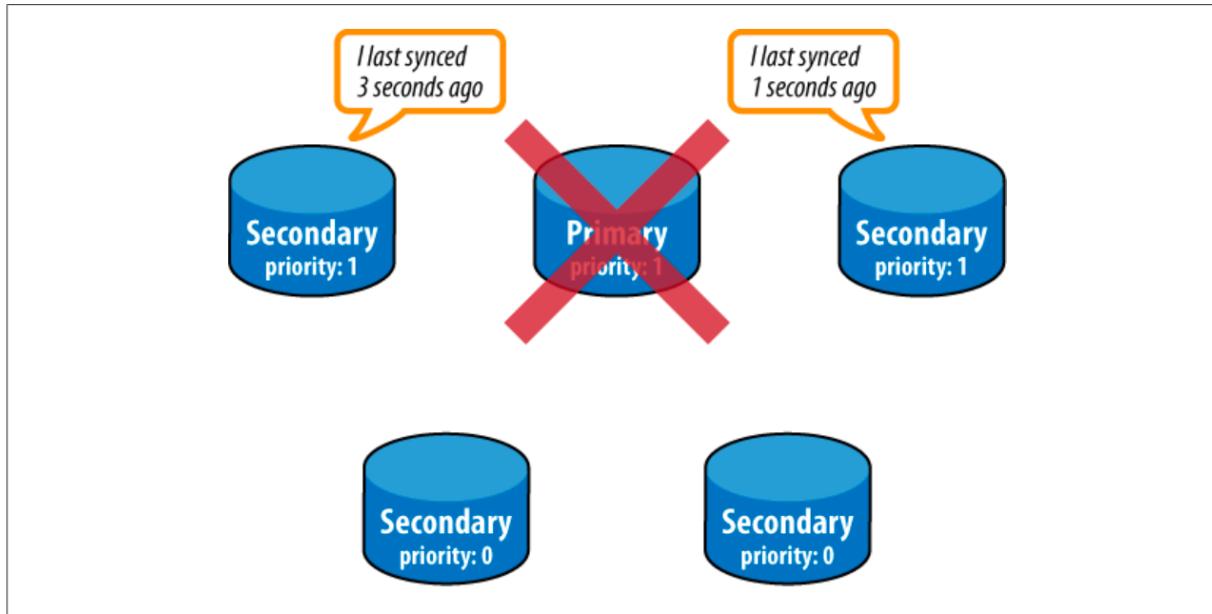


Figure 9-7. If the primary goes down, the highest-priority servers will compare how up-to-date they are

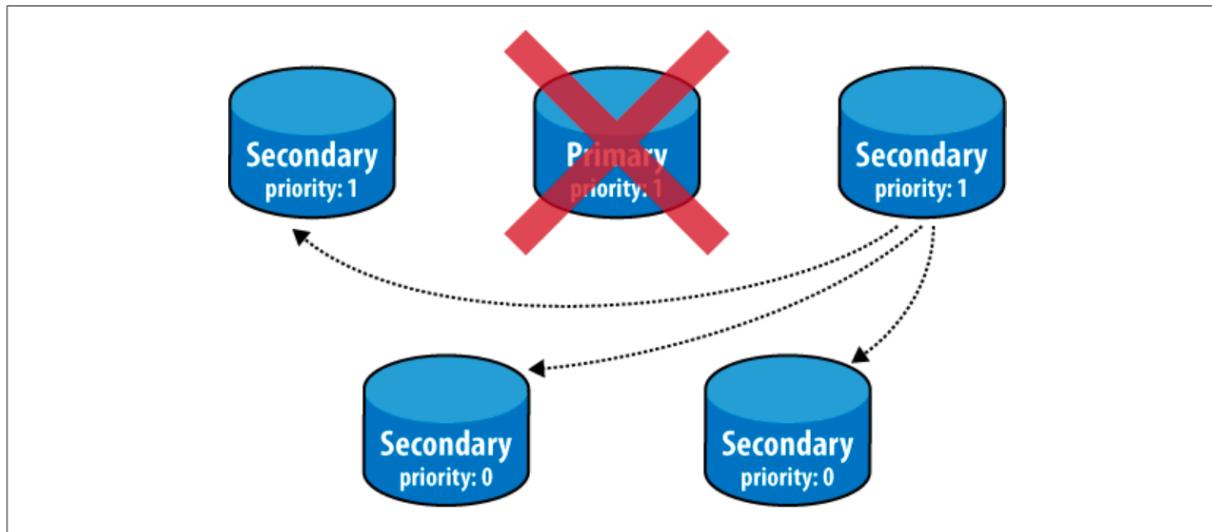
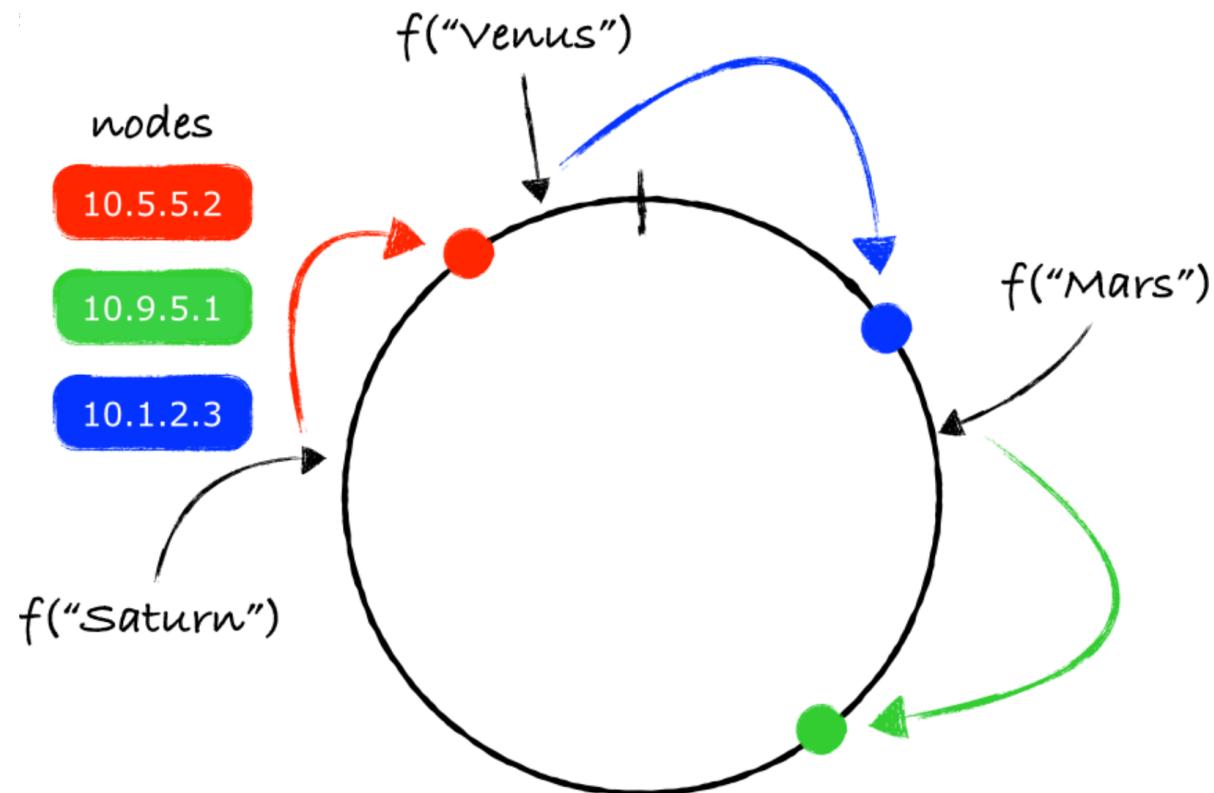


Figure 9-8. The highest-priority most-up-to-date server will become the new primary

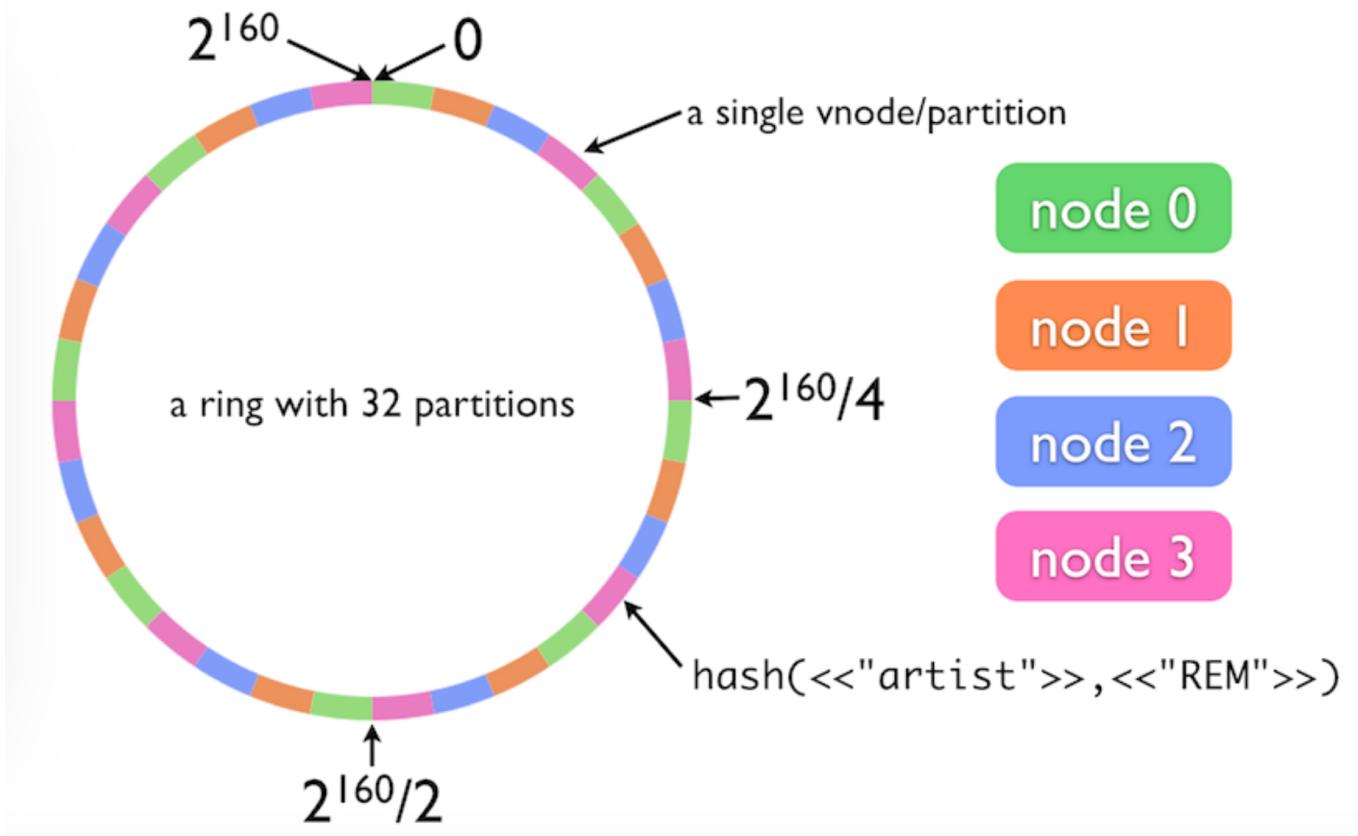
Figuras tomadas desde el libro → MongoDB: The definitive guide

## Consistent Hashing



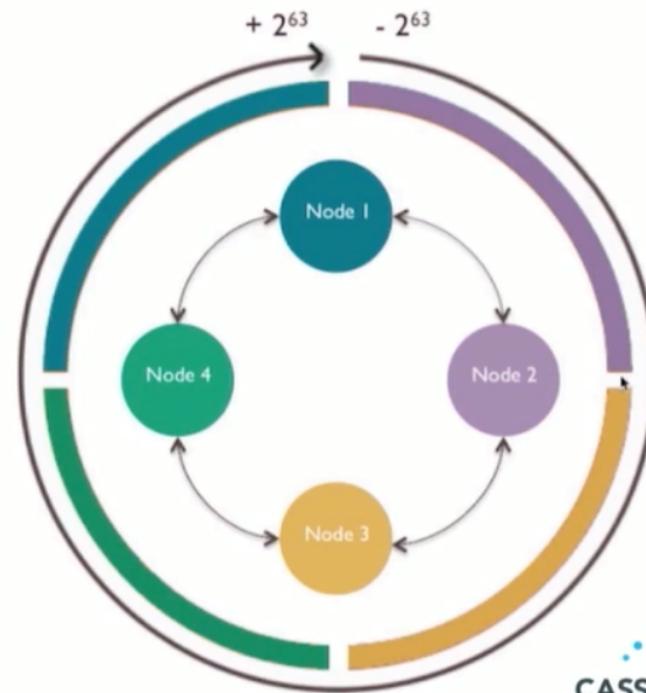
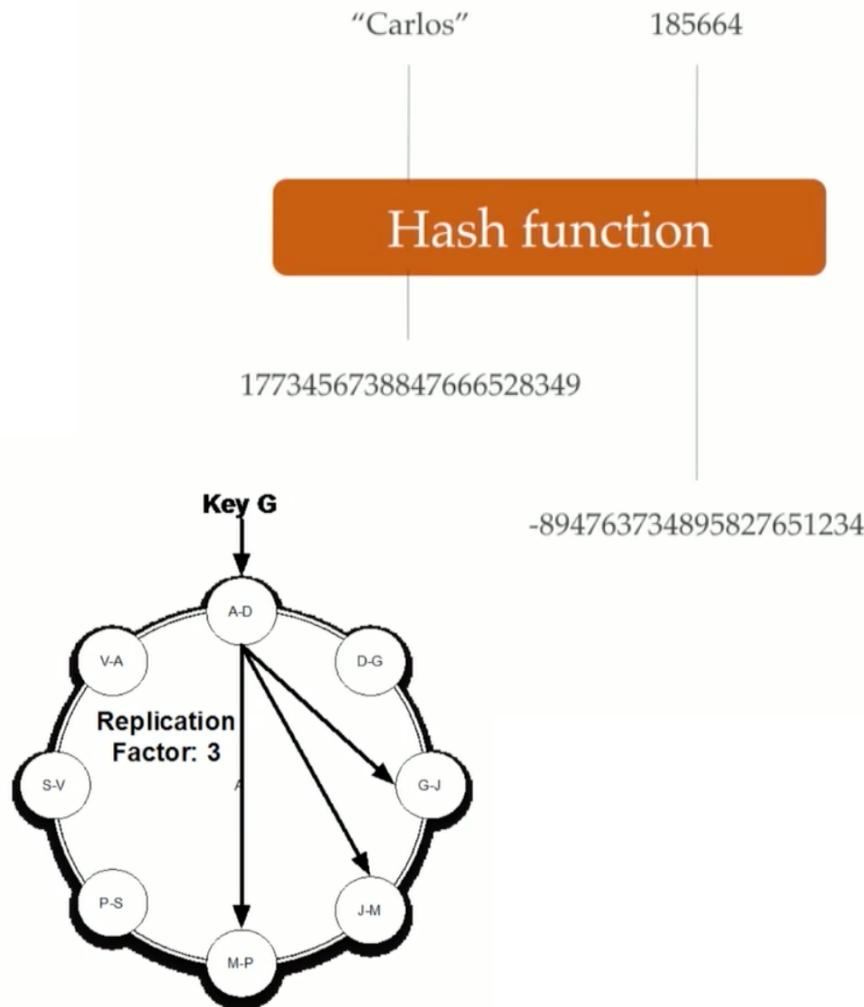
<https://akshatm.svbtle.com/consistent-hash-rings-theory-and-implementation>

## Balance de carga

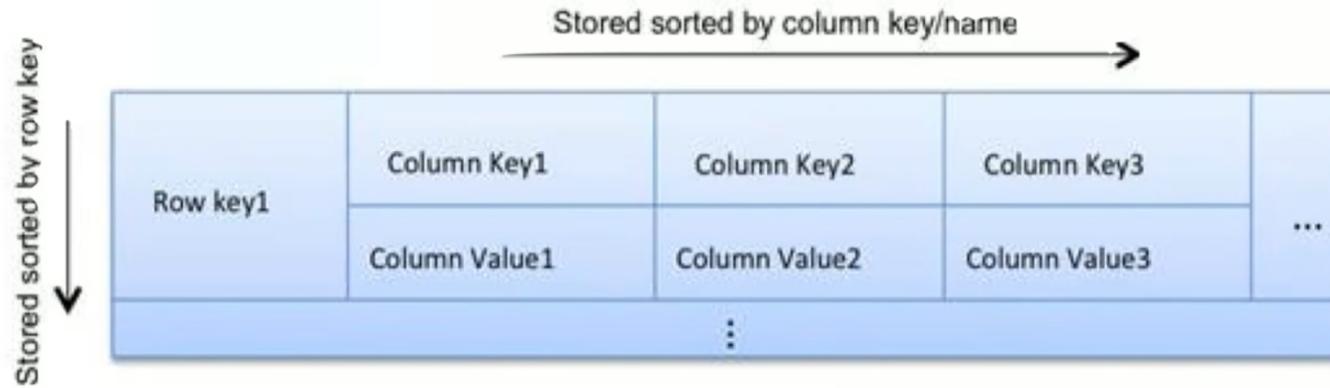


<https://docs.riak.com/riak/kv/2.2.3/learn/concepts/clusters/>

# Consistent Hashing



# Physical Data Layout

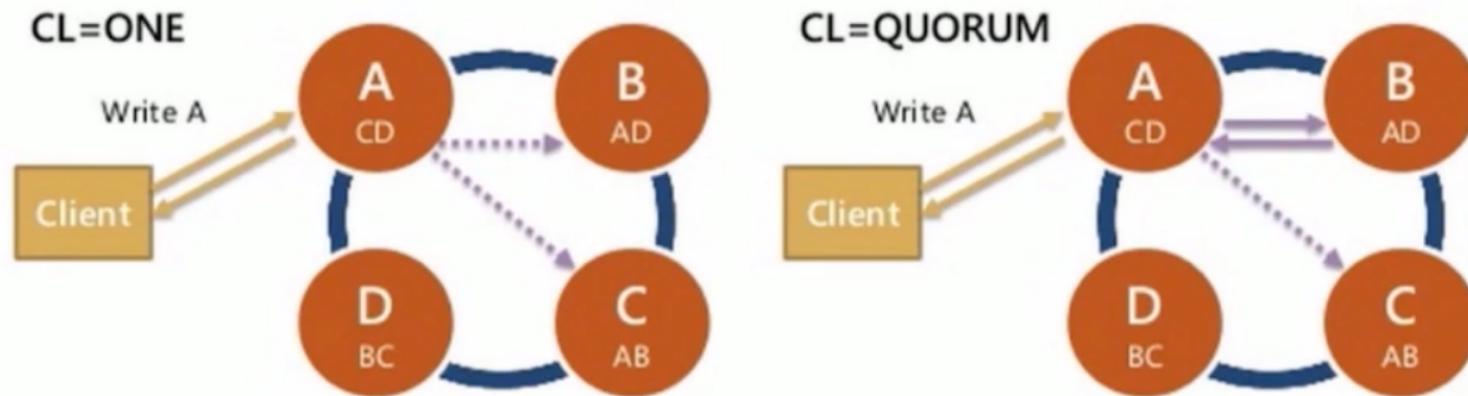
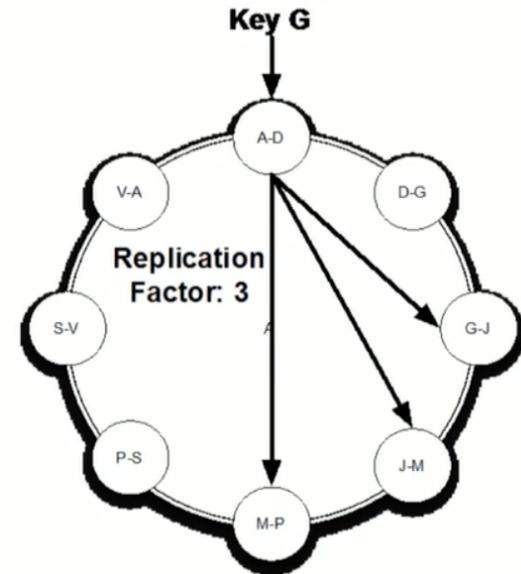


## Data organisation



# Consistency Level

How many replicas of your data must acknowledge?



# A complete read/write example

- RF = 3
- CL = QUORUM
- SELECT \* ... WHERE id = f81d4fae-...

